

I n s i d e M a c O S X

Kernel Environment



Preliminary 4/11/00
Technical Publications
©2000 Apple Computer, Inc.



Apple Computer, Inc.

© 2000 Apple Computer, Inc.

All rights reserved.

No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means, mechanical, electronic, photocopying, recording, or otherwise, without prior written permission of Apple Computer, Inc., with the following exceptions: Any person is hereby authorized to store documentation on a single computer for personal use only and to print copies of documentation for personal use provided that the documentation contains Apple's copyright notice.

The Apple logo is a trademark of Apple Computer, Inc.

Use of the "keyboard" Apple logo (Option-Shift-K) for commercial purposes without the prior written consent of Apple may constitute trademark infringement and unfair competition in violation of federal and state laws.

No licenses, express or implied, are granted with respect to any of the technology described in this book. Apple retains all intellectual property rights associated with the technology described in this book. This book is intended to assist application developers to develop applications only for Apple-labeled or Apple-licensed computers.

Every effort has been made to ensure that the information in this document is accurate. Apple is not responsible for typographical errors.

Apple Computer, Inc.

1 Infinite Loop

Cupertino, CA 95014

408-996-1010

Apple, the Apple logo, and Macintosh are trademarks of Apple Computer, Inc., registered in the United States and other countries.

Simultaneously published in the United States and Canada.

Even though Apple has reviewed this manual, **APPLE MAKES NO WARRANTY OR REPRESENTATION, EITHER EXPRESS OR IMPLIED, WITH RESPECT TO THIS MANUAL, ITS QUALITY, ACCURACY, MERCHANTABILITY, OR FITNESS FOR A PARTICULAR PURPOSE. AS A RESULT, THIS MANUAL IS SOLD "AS IS," AND YOU, THE PURCHASER, ARE ASSUMING THE ENTIRE RISK AS TO ITS QUALITY AND ACCURACY.**

IN NO EVENT WILL APPLE BE LIABLE FOR DIRECT, INDIRECT, SPECIAL, INCIDENTAL, OR CONSEQUENTIAL DAMAGES RESULTING FROM ANY DEFECT OR INACCURACY IN THIS MANUAL, even if advised of the possibility of such damages.

THE WARRANTY AND REMEDIES SET FORTH ABOVE ARE EXCLUSIVE AND IN LIEU OF ALL OTHERS, ORAL OR WRITTEN, EXPRESS OR IMPLIED. No Apple dealer, agent, or employee is authorized to make any modification, extension, or addition to this warranty.

Some states do not allow the exclusion or limitation of implied warranties or liability for incidental or consequential damages, so the above limitation or exclusion may not apply to you. This warranty gives you specific legal rights, and you may also have other rights which vary from state to state.

1 About This Book

The purpose of this book is to provide fundamental high-level information about the Mac OS X core operating system architecture. It also provides background for system programmers and developers of device drivers, file systems, and network extensions. The book concentrates on those areas where Mac OS X system architecture differs from other, similar operating systems.

Therefore, this book does not delve deeply into the specific APIs or programmatic use of the individual components of the Mac OS X core operating system, collectively known as the **kernel**. These components include **Mach**, **BSD**, the I/O Kit, networking, and the file system. To learn more about how to program for these components, you should see the specific documentation for each of them. The succeeding chapters of this book will provide pointers to more in-depth reference material.

Audience Profile

This book has a wide and diverse audience – specifically, the set of potential system software developers for Mac OS X. It is anticipated that the audience will consist primarily of the following sorts of developers:

- **Device driver writers** – Device driver writers make up the largest portion of the audience. You will be some of the first developers to start writing code for Mac OS X. Most device driver writers will have come from one of the following platform backgrounds:
 - **Generic device drivers (third party solutions)**—Your company writes drivers for many operating systems – Mac OS, Microsoft Windows, Linux,

and so forth. You'd like to know how writing for Mac OS X will be different (and how it will be similar).

- **UNIX platforms and variants of UNIX platforms such as FreeBSD, Linux, Solaris, and others**—You've been writing drivers for platforms such as Linux and FreeBSD. You want to know how to modify your code (or change your habits) when writing for Mac OS X. You may have certain preconceived notions about writing device drivers. The kernel environment model in Mac OS X differs in several respects from what you are used to; you'll need to understand those differences.
- **Windows NT**—You have been writing for the Windows NT platform. Now, you have decided to broaden your scope. You need to know how to write for Mac OS X.
- **Mac OS (Classic)** —You have been writing drivers for Mac OS for a long time and you know everything there is to know about Mac OS 8 and 9. However, Mac OS X is different. You need to know how to modify your code (or change your habits) when porting to Mac OS X.
- **NeXT (OpenStep)** —You wrote drivers for OpenStep. You're hoping that some of what you learned is still applicable. You need to know how to modify your code (or change your habits) when writing for Mac OS X.
- **Network extension writers** —You need to know how the networking subsystem fits in with the rest of the core operating system. You come from a platform background similar to the device driver writers.
- **File system writers** —You want to support a file system such as AFS or NTFS. You need to understand how to fit your code into Mac OS X.
- **Developers of software requiring very low level access to file system data** — You are writing software that needs low level access to the file system, applications such as “on-the-fly” compression, encryption and virus checking. You need to understand how to write Virtual File System stacks to add value on top of Mac OS X.
- **System programmers familiar with BSD, Linux, and similar operating systems** — As a system programmer, you're wondering what Mac OS X has to offer you. This book will address the differences between Mac OS X and the “standard” BSD and Mach 3.0 implementations.
- **Customers with special requirements** — Because the Mac OS X kernel technology is Open Source, some developers will be planning to make changes to the underlying operating system in order to meet special requirements at

their site. For example, a University customer may need to add Kerberos support. This book will tell you how the parts of the Mac OS X kernel fit together and interact.

- **Applications developers, students and others** – You’re not a system programmer, but you’re interested in how Mac OS X is put together. You may already be familiar with BSD, Linux, or other UNIX variants and possibly Windows NT as well. Although you don’t expect to need to know a great deal about the kernel environment, you are nonetheless interested in some details of memory allocation, process management, and the like.

Road Map

The goal of this book is to describe the underlying global concepts of the core operating system development environment. That is, it describes shared concepts that are not specific to any one of the primary subsystems: Mach, the I/O Kit, POSIX/BSD, file systems, or networking. All concepts should be applicable to each of these subsystems and are therefore useful to developers in all “camps” (such as device-driver writers).

This book will not delve deeply into the specific APIs or programmatic use of the individual subsystems of the operating system. Each of these subsystems will be the subject of its own documentation.

The chapters of this book describe the kernel environment from different angles. Discussion of specific APIs, however, will be left to more in-depth component-specific documentation.

The next chapter provides an overview of Mac OS X architecture. There follow several chapters that discuss each of the architectural components of Mac OS X in more detail, one chapter per component. These are followed in turn by a chapter that discusses extending the kernel, from a conceptual viewpoint. The last chapter covers available kernel services. For each service, it provides a brief description as well as listing which components are either a provider or a client.

The book ends with a glossary of terms used throughout the preceding chapters and a comprehensive reference bibliography. The goal of this book is very broad — to provide a firm grounding in the fundamentals of Mac OS X system

About This Book

programming for developers from many backgrounds. Unfortunately, to do a complete and comprehensive job would fill an entire library, rather than one book.

Instead, this book includes numerous pointers and references to books and other publications already in existence. Some of these are Apple publications; others are external documents. To make things easier, the bibliographic references are grouped into categories.

The glossary covers many of the terms used throughout the earlier chapters of this book; these terms will be highlighted in bold when first used. Rather than stop and define each term as it appears, the definitions have been moved to the end. If a term seems familiar, it probably means what you think it does. If it's unfamiliar, check the glossary. In any case, all readers may want to skim through the glossary, in case there are subtle differences between Mac OS X usage and other operating systems.

By the time you have finished this book, you should have a basic understanding of Mac OS X system internals and how to begin programming Mac OS X system software. You should also have a good idea of what you'll need to read next.

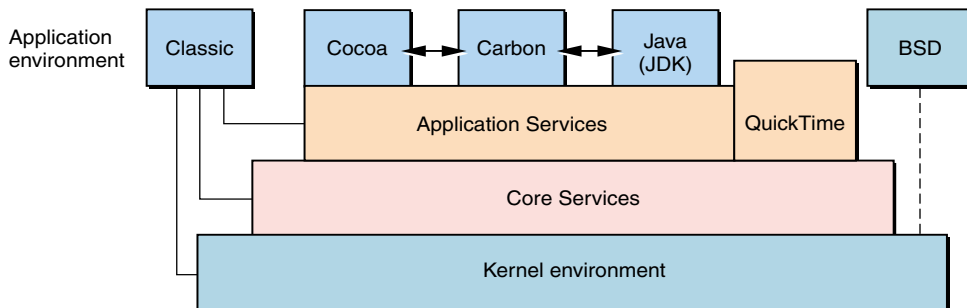
2 Kernel Architecture

Mac OS X provides many benefits to the Macintosh user and developer communities. These benefits include improved reliability and performance, enhanced networking features, an object-based system programming interface, and increased support for industry standards.

In creating Mac OS X, Apple has completely re-engineered the Mac OS core operating system. Forming the foundation of Mac OS X is the kernel. The figure below illustrates the Mac OS X architecture.

Figure 2-1

Mac OS X architecture



The kernel provides many enhancements for Mac OS X. These include **preemption**, **memory protection**, enhanced performance, improved networking facilities, support for both Macintosh (Extended and Standard) and non-Macintosh (UFS, ISO 9660) file systems, object-oriented APIs, and more. Two of these features, preemption and memory protection, lead to a more robust environment.

In Mac OS 8 and 9, applications cooperate to share processor time. Similarly, all applications share the memory of the computer among them. Mac OS 8 and 9 are **cooperative multitasking** environments. Everyone is hurt if a single application doesn't cooperate. On the other hand, real-time applications such as multimedia need to be assured of predictable, time-critical, behavior.

In contrast, Mac OS X is a **preemptive multitasking** environment. In Mac OS X, the kernel provides enforcement of cooperation, scheduling processes to share time (preemption). This supports real time behavior in applications that require it.

In Mac OS X, processes do not normally share memory. Instead, the kernel assigns each **process** its own **address space**, controlling access to these address spaces. This control ensures that no application can inadvertently access or modify another application's memory (protection). Size is not an issue; with Mac OS X's virtual memory system, each application has access to its own 4 GB memory address space.

Viewed together, all applications are said to run in "user space", but this does not imply that they share memory. User space is simply a term for the combined address spaces of all user-level applications. The kernel itself has its own address space, called kernel space. In Mac OS X, no application can modify the memory of the system software (the kernel).

Although user processes do not share memory as in Mac OS 8 and 9, communication between applications is still possible. For example, the kernel offers a rich set of primitives to permit some sharing of information among processes. These primitives include shared libraries and frameworks. Mach messaging provides another approach, handing memory from one process to another. Unlike Mac OS 8 and 9, however, memory sharing cannot occur without explicit action by the programmer.

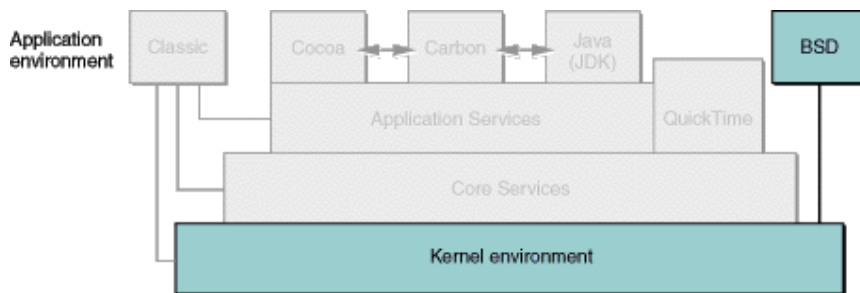
Darwin OS

The Mac OS X kernel is a key part of Apple's **Open Source** initiative. The Mac OS X kernel is also the core of an operating system product called **Darwin OS**. Darwin OS is a complete operating system based on many of the same technologies that underlie Mac OS X. However, Darwin OS does not include Apple's proprietary graphics or applications layers, such as Quartz, QuickTime, or OpenGL.

Kernel Architecture

Figure 1-2 below shows the relationship between Darwin OS and Mac OS X. Both build upon the same kernel, but Mac OS X adds Core services, Application services and QuickTime, as well as the **Carbon**, **Cocoa**, **Classic**, and Java (JDK) application environments. Both Darwin OS and Mac OS X include the BSD application environment (command line); however, in Mac OS X, this environment is usually hidden.

Figure 2-2 Darwin OS is a **subset of Mac OS X**



Darwin technology is based on **FreeBSD**, 4.2BSD Lite, Mach 3.0, and Apple technologies. Best of all, Darwin technology is Open Source technology, which means developers have full access to the source code. In effect, Mac OS X third-party developers can be part of the Darwin OS core system software development team. Developers can also see how Apple is doing things in the core operating system, adopting (or adapting) code to use within their own products. Refer to the **Apple Public Source License** for details.

Because the same system software forms the core of both Mac OS X and Darwin OS, system software developers will be able to write software that will run on both Mac OS X and Darwin OS with few, if any, required changes. The only difference might be in the way the software interacts with the application environment.

The Mac OS X core operating system is based on proven technology from many sources. A large portion of this technology is derived from FreeBSD, a version of 4.4BSD that offers advanced networking, performance, security and compatibility features. Other parts of the system software, such as Mach, are based on technology previously used in Apple's MkLinux project, in Mac OS X Server, and

in technology acquired from NeXT. Much of the code is platform-independent. All of the core operating system code is available in source form.

The core technologies have been chosen for several reasons. Mach provides a clean set of abstractions for dealing with memory management, inter-process (and interprocessor) communication, and other low-level operating system functions. In today's rapidly changing hardware environment, this provides a useful layer of insulation between the operating system and the underlying hardware.

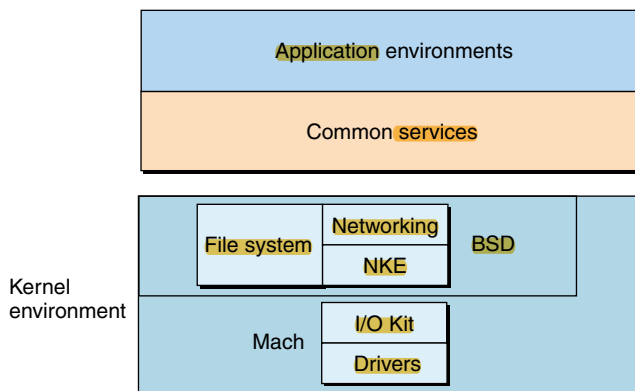
BSD is a carefully engineered, mature operating system with many capabilities. In fact, most of today's commercial Linux, UNIX, and other similar operating systems contain a great deal of BSD code. BSD also provides a set of industry-standard APIs.

New technologies, such as the I/O Kit and Network Kernel Extensions (NKE), have been designed and engineered by Apple to take advantage of advanced capabilities, such as those provided by an object-oriented programming model. Mac OS X combines these new technologies with time-tested industry standards to create an operating system that is stable, reliable, flexible, and programmable.

Architecture

The **foundation layer** of **Darwin OS and Mac OS X** is composed of several **architectural components**, as shown in the figure below. Taken together, these components form the **kernel environment** or simply, the kernel.

Figure 2-3 Mac OS X kernel architecture



Important

Note that Mac OS X uses the term **kernel** somewhat differently than you may be used to seeing it used.

A kernel, in traditional operating system terminology, is a **small nucleus of software** that provides only **the minimal facilities** necessary for **implementing additional operating-system services**.

(The Design and Implementation of the 4.4 BSD Operating System, McKusick, et. al.)

Instead, **Mac OS X** uses the term **kernel** to **refer to everything** that **executes in the kernel address space**.

The Mac OS X kernel includes Mach, BSD, the I/O Kit, file systems, and networking components. Each of these components is described briefly in the following sections. For further details, refer to the specific component chapters or to the reference material listed in the Bibliography.

Mach

Mach manages **processor resources** such as **CPU usage and memory**, handles **scheduling**, provides **memory protection**, and provides a **messaging centered**

infrastructure to the rest of the operating system layers. The Mach component provides:

- untyped **IPC and RPC**
- support for **SMP**
- support for **real-time services**
- external **pager**
- **modular architecture**
- improved **performance**

BSD

The BSD layer is based on 4.4BSD, specifically FreeBSD. **Above the Mach kernel**, this layer provides “OS personality” APIs and services. The BSD component provides:

- **file** systems
- **networking**
- basic **security policies** such as **user IDs and permissions**
- the **system framework**—a mechanism for **exporting APIs to the application layers**
- the BSD **process model**, including **process IDs and signals**
- FreeBSD **kernel APIs**
- **Pthreads** (POSIX threads implementation)

I/O Kit

The I/O Kit provides a framework for **simplified driver development**, supporting many categories of devices. The I/O Kit features an **object-oriented I/O architecture** implemented in Embedded C++, a subset of C++. The I/O Kit framework is both **modular and extensible**. The I/O Kit component provides:

- true **plug and play**
- **dynamic device** management

- **dynamic** (“on-demand”) **loading** of drivers
- **power management** for desktop systems as well as portables
- **multiprocessor capabilities**

Networking

Mac OS X networking takes advantage of **BSD’s advanced networking capabilities** to provide support for modern features, such as **Network Address Translation (NAT)** and **firewalls**. The networking component provides:

- 4.4BSD TCP/IP stack and socket APIs
- support for both IP and AppleTalk
- **multi-homing**
- routing
- **multicast** support
- server **tuning**
- **socket-based** AppleTalk
- Mac OS **Classic support**
- Carbonized Open Transport APIs

File Systems

Mac OS X provides support for **numerous types of file systems**, including **HFS**, **HFS+**, **UFS**, **NFS**, **ISO 9660**, and others. The **default** type is **HFS+**; Mac OS X **boots** (and “roots”) from **HFS+**. Advanced features of Mac OS X file systems include an enhanced Virtual File System (**VFS**) design. VFS provides for a layered architecture (file systems are **stackable**). The file system component provides:

- **UTF-8 (Unicode)** support
- increased **performance**

Kernel Extensions

Mac OS X provides a kernel extension mechanism as a means of **allowing dynamic loading of pieces of code into the kernel**, without the need to recompile. These pieces of code are known generically as **plug-ins** or in the Mac OS X kernel as **kernel extensions** or **KEXTs**.

Because KEXTs provide both modularity and dynamic loadability, they are a natural choice for any relatively self-contained service that requires access to kernel internal interfaces. Many of the components of the Core OS support this extension mechanism, although in different ways.

For example, some of the new networking features involve the use of network kernel extensions (**NKEs**). The ability to dynamically add a new file system implementation is based on VFS KEXTs. Device drivers and device families in the I/O Kit are implemented using KEXTs. For developers writing drivers, or code to support a new volume format or networking protocol, KEXTs make development much easier. KEXTs will be discussed in more detail in the chapter Extending the Kernel.

3 Mach

The fundamental services and primitives of the Mac OS X kernel are based on Mach 3.0. Apple has modified and extended Mach to better meet Mac OS X functional and performance goals.

Mach 3.0 was originally conceived as a simple, extensible, communications microkernel. It is capable of running standalone, with other traditional operating system services such as I/O, file systems, and networking stacks running as user-mode servers.

However, in Mac OS X, Mach is linked with other kernel components into a single kernel address space. This is primarily for performance; it is much faster to make a direct call between linked components than it is to send messages or do RPC between separate tasks. This modular structure results in a more robust and extensible system than a large, purely monolithic kernel would allow, without the performance penalty of a pure microkernel.

Thus in Mac OS X, Mach is not primarily a communication hub between clients and servers. Instead, its value consists of its abstractions, its extensibility, and its flexibility. In particular, Mach provides:

- Object-based APIs with communication channels (e.g., ports) as object references
- Highly parallel execution, including preemptively scheduled threads and support for SMP
- A flexible scheduling framework, with support for real-time usage
- A complete set of IPC primitives, including messaging, RPC, synchronization, and notification
- Support for large virtual address spaces, shared memory regions, and memory objects backed by persistent store

- Proven **extensibility** and **portability**, for example **across instruction set architectures** and in **distributed environments**
- **Security** and **resource management** as **a fundamental principle of design**; all **resources are virtualized**

Mach Kernel Abstractions

Mach provides **a small set of abstractions**, which have been designed to be both **simple** and **powerful**. The main kernel abstractions are

- **Task** — The unit of **resource** ownership, consisting of a **virtual address space**, a **port right namespace**, and a set of **threads**.
- **Thread** — The **unit of CPU execution**.
- **Address space** — In conjunction with **memory managers**, Mach implements the notion of a **sparse virtual address space** and **shared memory**.
- **Memory object** — **The internal unit of memory management**. Memory objects include **named entries and regions**; they are **representations of potentially-persistent data** which may be **mapped into address spaces**.
- **Port** — A secure, simplex **communication channel**, accessible only via **send/receive** capabilities (**rights**).
- **IPC** — **Message queues**, **remote procedure calls (RPC)**, **notifications**, **semaphores**, and **lock** sets.
- **Time** — **Clocks**, **timers**, and **waiting**.

At the **trap** level, the **interface to most Mach abstractions** consists of messages sent to and from kernel ports representing those objects. The trap-level interfaces (such as `mach_msg_overwrite_trap`) and **message formats** are themselves abstracted in normal usage by the **Mach Interface Generator (MIG)**. MIG is used to **compile procedural interfaces** to the **message-based APIs**, based on descriptions of those APIs.

Tasks And Threads

Mac OS X processes are implemented on top of Mach tasks, and POSIX threads (Pthreads) are implemented on top of Mach threads. A thread is Mach's notion of the point of control. A task exists to provide resources for its containing threads. This split is made to provide for parallelism and resource sharing.

A thread:

- Is a point of control flow in a task
- Has access to all of the elements of the containing task
- Potentially executes in parallel with other threads, even threads within the same task
- Has minimal state for low overhead

A task:

- Is a collection of system resources—these resources, with the exception of the address space, are referenced by ports. These resources may be shared with other tasks if rights to the ports are so distributed.
- Provides a large, potentially sparse address space, referenced by machine address; portions of this space may be shared through inheritance or external memory management.
- Contains some number of threads.

Note that a task has no life of its own; only threads execute instructions. When it is said “a task Y does X” what is really meant is that “a thread contained within task Y does X”.

A task is a fairly expensive entity. It exists to be a collection of resources. All of the threads in a task share everything. Two tasks share nothing without an explicit action (although the action is often simple) and some resources cannot be shared between two tasks at all (such as port receive rights).

A thread is a fairly lightweight entity. It is fairly cheap to create and has low overhead to operate. This is true because a thread has little state (mostly its register

state); its owning task bears the burden of resource management. On a multiprocessor it is possible for multiple threads in a task to execute in parallel. Even when parallelism is not the goal, multiple threads have an advantage in that each thread can use a synchronous programming style, instead of attempting asynchronous programming with a single thread attempting to provide multiple services.

A thread is the basic computational entity. A thread belongs to one and only one task that defines its virtual address space. To affect the structure of the address space, or to reference any resource other than the address space, the thread must execute a special trap instruction which causes the kernel to perform operations on behalf of the thread, or to send a message to some agent on behalf of the thread. In general, these traps manipulate resources associated with the task containing the thread. Requests can be made of the kernel to manipulate these entities: to create them, delete them and affect their state.

Mach provides a flexible framework for thread scheduling policies. Early versions of Mac OS X support both the timesharing and fixed-priority policies. A timesharing thread's priority is raised and lowered to balance its resource consumption against other timesharing threads. Fixed-priority threads execute for a certain quantum, and then are put at the end of the queue of threads of equal priority. Setting a fixed priority thread's quantum to 'infinity' effectively makes the thread run-till-block within its priority. High priority real-time threads are usually fixed priority.

Future versions of Mac OS X will have additional scheduling policies, for more sophisticated real-time support.

Ports, Port Rights, Port Sets, and Port Name Spaces

With the exception of the task's virtual address space, all other system resources are accessed through a level of indirection known as a port. A port is an endpoint of a unidirectional communication channel between a client who requests a service and a server who provides the service. If a reply is to be provided to such a service request, a second port must be used.

In most cases, the resource that is accessed by the port (that is, named by it) is referred to as an object. Most objects named by a port have a single receiver and

potentially multiple senders. That is, there is exactly one receive port, and at least one sending port, for a typical object such as a message queue. The service to be provided by an object is determined by the manager that receives the request sent to the object. It follows that the receiver for ports associated with kernel-provided objects is the kernel and the receiver for ports associated with task-provided objects is the task providing that object.

For ports that name task-provided objects, it is possible to change the receiver of requests for that port to be a different task, for example by passing the port to that task in a message. A single task may have multiple ports that refer to resources it supports. For that matter, any given entity can have multiple ports that represent it, each implying different sets of permissible operations. For example, many objects have a **name port** and a **control port** (sometimes called the privileged port). Access to the control port allows the object to be manipulated; access to the name port simply names the object, for example, to return information about it.

Tasks have permissions to access ports in certain ways (send, receive, send-once); these are called **port rights**. A port can only be accessed via a right. Ports are often used to grant clients access to objects within Mach. Having the right to "send" to the object's IPC port denotes the right to manipulate the object in prescribed ways. As such, port right ownership is the fundamental security mechanism within Mach. Having a right to an object is to have a capability to access or manipulate that object.

Port rights can be copied and moved between tasks via IPC. Doing so, in effect, passes capabilities to some object or server.

One type of object referred to by a port is a **port set**. As the name suggests, a port set is a set of port rights which can be treated as a single unit when receiving a message or event from any of the members of the set. Port sets permit one thread to wait on a number of message and event sources, for example in **work loops**.

Traditionally in Mach, the communication channel denoted by a port was always a queue of **messages**. However, Mac OS X supports additional types of communication channels, and these new types of IPC object are also represented by ports and port rights. See the section, [Task To Task Communication \(IPC\)](#), for more details about messages and other IPC types.

Ports and port rights do not have system-wide names that allow arbitrary ports or rights to be manipulated directly. Ports can only be manipulated by a task if it has a port right in its **port name space**. A port right is specified by a **port name**, which is

an integer index into a 32-bit port name space. Each task has associated with it a single port name space.

Tasks acquire port rights when another task explicitly inserts them into its name space, when they receive rights in messages, by creating objects which return a right to the object, and via Mach calls for certain special ports (`mach_thread_self`, `mach_task_self`, and `mach_reply_port`.)

Memory Management

As with most modern operating systems, Mach provides addressing to large, sparse, virtual address spaces. Runtime access is made via virtual addresses which may not correspond to locations in physical memory at the initial time of the attempted access. Mach is responsible for reconciling a requested access in virtual space with a location in physical memory. It does so through demand paging.

A range of a virtual address space is populated with data when a memory object is mapped into that range. All data in an address space is ultimately provided through memory objects. Mach will ask the owner of a memory object (a pager) for the contents of a page when establishing it in physical memory, and will return the possibly-modified data to the pager before reclaiming the page. Mac OS X includes two built-in pagers — the default pager and the vnode pager.

The default pager handles non-persistent or anonymous memory. Anonymous memory is zero-initialized, and it exists only during the life of a task. The vnode pager maps files into memory objects. Mach exports an interface to memory objects to allow their contents to be contributed by user-mode tasks. This interface is known as the External Memory Management Interface, or EMMI.

The memory management subsystem exports virtual memory handles known as named memory entries. Like most kernel resources, these are denoted by ports. Having a named memory entry handle allows the owner to map the underlying virtual memory object or to pass the right to map the underlying object to others. Mapping a named entry in two different tasks results in a shared memory window between the two tasks, thus providing a flexible method for establishing shared memory.



Address ranges of virtual memory space may also be populated through direct allocation (using `vm_allocate`). The underlying virtual memory object is anonymous and backed by the default pager. Shared ranges of an address space may also be set up via inheritance. When new tasks are created they are cloned from a parent. This cloning pertains to the underlying memory address space as well. Mapped portions of objects may be inherited as a copy, or as shared, or not at all, based on attributes associated with the mappings. Mach practices a form of delayed copy known as **copy-on-write** to optimize the performance of inherited copies on task creation.

Rather than directly copying the range, a copy-on-write optimization is accomplished by protected sharing. The two tasks both share the memory to be copied, but with read-only access. When either task attempts to modify a portion of the range, that portion is copied at that time. This lazy evaluation of memory copies is an important optimization that permits simplifications in several areas, notably the messaging APIs.

One other form of sharing is provided by Mach, through the export of **named regions**. A named region is a form of a named entry, but instead of being backed by virtual memory object, it is backed by a virtual map fragment. This fragment may hold mappings to numerous virtual memory objects. It is mappable into other virtual maps, providing a way of inheriting not only a group of virtual memory objects but also their existing mapping relationships. This feature offers significant optimization in task setup, for example when sharing a complex region of the address space used for shared libraries.

Task To Task Communication (IPC)

Communication between tasks is an important element of the Mach philosophy. Mach supports a client/server system structure in which tasks (clients) access services by making requests of other tasks (servers) via messages sent over a communication channel.

The endpoints of these communication channels in Mach are called ports, while port rights denote permission to use the channel. The forms of IPC provided by Mach include

- Message queues

- Semaphores
- Notifications
- Lock sets
- Remote procedure calls (RPC)

The **type** of IPC object denoted by the port **determines** the **operations permissible on that port**, and how (and whether) data transfer occurs.

Important

The IPC facilities in Mac OS X are in a state of **transition**. In early versions of the system, not all of these IPC types may be implemented.

There are two fundamentally different **Mach APIs** for **raw manipulation of ports** — the **mach_msg** family, and the **mach_ipc** family of calls. Within reason, both families may be **used with any IPC object**; however, the **mach_ipc** calls are **preferred in new code**. The **mach_msg** calls are supported for **legacy code but deprecated**; they are **stateless**. The **mach_ipc** calls are **stateful where appropriate in order to support the notion of a transaction**.

IPC Transactions and Event Dispatching

When a thread calls **mach_ipc_dispatch**, it will **repeatedly process events coming in on the registered portset**. These events could be **an argument block** from an RPC object (as the results of a client's call), **a lock object** being taken (as a result of some other thread's releasing the lock), **a notification or semaphore** being posted, or a **message** coming in from a traditional message queue.

These events are handled via callouts from **mach_msg_dispatch**. Some events imply a transaction during the lifetime of the callout, because they have state. In the case of a lock, the **state** is the **ownership of the lock**. When the callout **returns**, the lock is **released**. In the case of RPC, the state is the client's identity, the argument block, and the reply port. When the callout returns, the reply is sent.

When the callout **returns**, the **transaction** (if any) is **completed**, and the **thread waits for the next event**. The **mach_ipc_dispatch** facility is intended to **support work loops**.

Message Queues

Originally, the sole style of IPC in Mach was the message queue. Only one task can hold the receive right for a port denoting a message queue. This one task is allowed to receive (read) messages from the port queue. Multiple tasks can hold rights to the port that allow them to send (write) messages into the queue.

A task communicates with another task by building a data structure that contains a set of data elements, and then performing a message-send operation on a port for which it holds send rights. At some later time, the task with receive rights to that port will perform a message-receive operation.

A message may consist of some or all of the following:

- Pure data
- Copies of memory ranges
- Port rights
- Kernel implicit attributes, such as the sender's security token

The message transfer is an asynchronous operation. The message is logically copied into the receiving task, possibly with copy-on-write optimizations. Multiple threads within the receiving task can be attempting to receive messages from a given port, but only one thread will receive any given message.

Semaphores

Semaphore IPC objects support “wait”, “post”, and “post all” operations. These are counting semaphores, in that posts are saved (counted) if there are no current waiting threads. “Post all” wakes up all currently waiting threads. There is no data associated with a semaphore.

Notifications

Like semaphores, notification objects also support the “post” and “wait” methods, but with the addition of a state field. The state is a fixed size, fixed format field which is defined when the notification object is created. Each post updates the state field; there is a single state, which is overwritten by each post.

Lock

A lock is a **mutex**. The primary interfaces to locks are transaction oriented (see the section on IPC transactions, above). During the transaction, the thread holds the lock. When it returns from the transaction, the lock is released. There is no data associated with the lock.

Remote Procedure Calls (RPC)

As the name implies, an RPC object is designed to facilitate and optimize remote procedure calls. The primary interfaces to RPC objects are transaction oriented (see IPC Transactions and Event Dispatching, above.)

When an RPC object is created, a set of argument block formats is defined. When an RPC call (a send on the object) is made by a client, it causes a message in one of the pre-defined formats to be created and queued on the object, then eventually passed to the server (the receiver). When the server returns from the transaction, the reply is returned to the sender. Mach tries to optimize the transaction by executing the server using the client's resources; this is called **thread migration**.

Time Management

The traditional abstraction of time in Mach is the **clock**, which provides a set of asynchronous alarm services based on `mach_timespec_t`. There are one or more clock objects, each defining a monotonically increasing time value expressed in nanoseconds. The real-time clock is built in, and is the most important, but there may be other clocks for other notions of time in the system. Clocks support operations to get the current time, sleep for a given period, set an alarm (a notification which is sent at a given time), and so forth.

The `mach_timespec_t` APIs are deprecated in Mac OS X. The newer and preferred APIs are based on timer objects, which in turn use `AbsoluteTime` as the basic data type. `AbsoluteTime` is a machine-dependent type, typically based on the platform-native time base. Routines are provided to convert `AbsoluteTime` to and from other data types, such as nanoseconds. Timer objects support asynchronous,

M a c h

drift-free notification, cancellation, and premature alarms. They are more efficient and permit higher resolution than clocks.

Important

As with several other Mach services, time management is in a state of transition in Mac OS X. Early versions of the system may not implement timer objects.

M a c h

4 BSD

BSD Features

The BSD portion of the Mac OS X kernel is derived from FreeBSD, a version of 4.4BSD that offers advanced networking, performance, security and compatibility features. Specifically, the BSD layer is based on the 4.4BSD-Lite2 release from Computer Systems Research Group (CSRG) at the University of California at Berkeley. BSD provides many advanced features, including

- Preemptive multitasking with dynamic priority adjustment—Smooth and fair sharing of the computer between applications and users is ensured, even under the heaviest of loads.
- Multiuser access—Many people can use a Mac OS X system simultaneously for a variety of things. This means, for example, that system peripherals such as printers and tape drives are properly shared between all users on the system or the network and that individual resource limits can be placed on users or groups of users, protecting critical system resources from over-use.
- Strong TCP/IP networking with support for industry standards such as SLIP, PPP, NFS, DHCP and NIS—Mac OS X can inter-operate easily with other systems as well as act as an enterprise server, providing vital functions such as NFS (remote file access) and email services, or Internet services such as http, ftp, routing and firewall (security) services.
- Memory protection — Applications cannot interfere with each other. One application crashing will not affect others in any way.

- Virtual memory and dynamic memory allocation—Applications with large appetites for memory are satisfied while still maintaining interactive response to users. With Mac OS X's virtual memory system, each application has access to its own 4 GB memory address space; this should satisfy even the most memory-hungry applications.
- SMP support—Support is included for machines with multiple CPUs.
- Source code—Developers gain the greatest degree of control over the BSD programming environment because source is included.
- Support for kernel threads, based on Mach threads—User-level threading packages are implemented on top of kernel threads. Each kernel thread is an independently scheduled entity. When a thread from a user process blocks in a system call, other threads from the same process can continue to execute on that or other processors. By default, a process in the conventional sense has one thread, the **main thread**. A user process can use the POSIX thread API to create other user threads.

BSD Facilities

The facilities available to a user process are logically divided into two parts: kernel facilities directly implemented by code running in the operating system, and system facilities implemented either by the system, or in cooperation with a server process.

The facilities implemented in the kernel define the **virtual machine** in which each process runs. Like many real machines, this virtual machine has memory management hardware, an interrupt facility, timers and counters.

The virtual machine also allows access to files and other objects through a set of descriptors. Each descriptor resembles a device controller, and supports a set of operations. Like devices on real machines, some of which are internal to the machine and some of which are external, parts of the descriptor machinery are built into the operating system, while other parts are often implemented in server processes.

The kernel facilities available from the BSD component include

■ **Processes and protection —**

- ☐ Host and process identifiers
- ☐ Process creation and termination
- ☐ User and group IDs
- ☐ Process groups

■ **Memory management —**

- ☐ Text, Data, Stack, and Dynamic Shared Libraries
- ☐ Mapping pages
- ☐ Page protection control
- ☐ Synchronization primitives

■ **Signals —**

- ☐ Signal types
- ☐ Signal handlers
- ☐ Sending signals

■ **Timing and statistics —**

- ☐ Real time
- ☐ Interval time

■ **Descriptors —**

- ☐ Files
- ☐ Pipes
- ☐ Sockets
- ☐ POSIX shared memory
- ☐ POSIX semaphores

■ **Resource controls —**

- ☐ Process priorities
- ☐ Resource utilization and resource limits

- Quotas

■ **System operation support —**

- Bootstrap operations
- Shutdown operations
- Accounting

BSD system facilities (facilities that may interact with user space) include

- Generic input/output operations such as read and write, non-blocking and asynchronous operations
- File system operations
- Interprocess communication
- Handling of terminals and other devices
- Process control
- Networking operations

Differences between Mac OS X and 4.4BSD

Although the BSD portion of Mac OS X is derived from FreeBSD and 4.4BSDLite2, some changes have been made.

- The `sbrk()` system call for memory management has not been implemented in Mac OS X.
- A runtime model supporting dynamic shared libraries is included. This model uses **Mach-O** and **PEF** binary file formats, as well as **dyld** (dynamic link editor) and **CFM** (Code Fragment Manager) that use these formats respectively. The kernel supports `execve()` of Mach-O binaries. Mapping and management of Mach-O dynamic shared libraries, as well as launching of PEF-based applications, are performed by user-space code.
- Mac OS X does not support memory mapped devices through the `mmap()` API
- The `swapon()` call is not supported; `macx_swapon()` is the equivalent call from the Mach pager.

- The Unified Buffer Cache implementation in Mac OS X differs from that found in FreeBSD.

In addition, several new features have been added that are specific to the Mac OS X (Darwin) implementation of BSD. These features are not found in FreeBSD.

- Enhancements to file system buffer cache and file I/O clustering
 - Adaptive and speculative read ahead
 - User process controlled read ahead
 - Time aging of the file system buffer cache
- Enhancements to file system support
 - Implementation of Apple extensions for ISO-9660 file system
 - Multi-threaded asynchronous I/O for NFS
 - Addition of system calls to support semantics of Mac OS Extended file system
 - Additions to naming conventions for pathnames, as required for accessing multiple forks in Mac OS Extended file systems

For further reading

The BSD component of the Mac OS X kernel is complex. A complete description is beyond the scope of this document. However, many excellent references exist for this component. For further information, developers interested in BSD should be sure to refer to the Bibliography.

Although the BSD layer of Mac OS X is derived from 4.4BSD, keep in mind that it is not identical to 4.4BSD. Some functionality of 4.4 BSD has not been included in Mac OS X. Some new functionality has been added. The cited reference materials are recommended for additional reading. However, they should **not** be presumed as forming a definitive description of Mac OS X.

B S D

5 Device Drivers and the I/O Kit

In creating Mac OS X, Apple has completely redesigned the Macintosh I/O architecture. Developers who are already familiar with writing device drivers for Mac OS 8 and 9 or for BSD will discover that writing drivers for Mac OS X requires some new ways of thinking. However, once you are familiar with the new model, you should find that the I/O Kit makes writing device drivers easier and more efficient than ever before.

The I/O Kit uses an object-oriented programming model, implemented in **Embedded C++**. Use of object-oriented frameworks can dramatically increase developer productivity.

From a programming perspective, the I/O Kit provides an abstract view of the system hardware to the upper layers of Mac OS X. By starting with properly designed base classes, a developer gains a “head start” in writing a new driver; with much of the driver code already written, the developer needs only to fill in the specific code that makes their driver different.

Another part of the philosophy of the I/O Kit is to make the design completely open. Rather than hiding APIs in an attempt to protect developers from themselves, all of the I/O Kit source is available as part of Darwin OS. Even private classes are visible to the programmer. (However, as with all C++ programming, use of private classes is not recommended!)

Instead of hiding the interfaces, Apple’s designers have chosen to lead by example. Sample code and classes show the recommended (easy) way to write a driver. However, developers are not prevented from doing things the hard way (or the wrong way!). Instead, attention has been concentrated on making the “best” ways easy to follow.

Redesigning the I/O Model

Many developers may ask why Apple chose to redesign the I/O model. At first glance, it might seem that re-using the model from Mac OS 9 or FreeBSD would have been an easier choice. There are several reasons for the decision, however.

Neither the Mac OS 8 and 9 driver model nor the FreeBSD model offered a sufficient feature set to meet the needs of Mac OS X. The underlying operating system technology of Mac OS X is very different from that of Mac OS 8 and 9. The Mac OS X kernel is significantly more advanced than the previous Mac OS system architecture; Mac OS X needs to handle memory protection, preemption, multi-processing, and other features not present in previous versions of Mac OS. Although FreeBSD is capable of handling these features, the BSD model did not offer the automatic configuration, stacking, power management, or dynamic device loading features required in a modern, consumer-oriented operating system.

By redesigning the I/O architecture, Apple's engineers can take best advantage of the operating system features in Mac OS X. For example, virtual memory (VM) is not a fundamental part of the operating system in Mac OS 8 and 9. Thus, every driver writer must know about (and write for) VM. This has presented certain complications for developers. In contrast, Mac OS X has simplified driver interaction with VM. VM capability is inherent in the Mac OS X operating system and cannot be turned off by the user. Thus, VM capabilities can be abstracted into the I/O Kit and the code for handling VM need not be written for every driver.

Mac OS X offers developers an unprecedented opportunity to take advantage of hardware complexity without the requirement of encoding software complexity into each new device driver. Under Mac OS 9, for example, all software development kits (SDKs) were independent of each other, duplicating functionality between them. In Mac OS X, the I/O Kit is delivered as part of the single kernel development kit (KDK); all portions of the KDK rely on common underpinnings.

I/O Kit Architecture

A device driver provides an abstract view of the hardware it is driving. These abstractions are defined in software by the I/O Kit **families** and **nubs**. Drivers provide the implementation behind the abstraction.

A family defines a collection of software abstractions that are common to all devices (instances) of a particular category (device class). Families provide functionality and services to drivers.

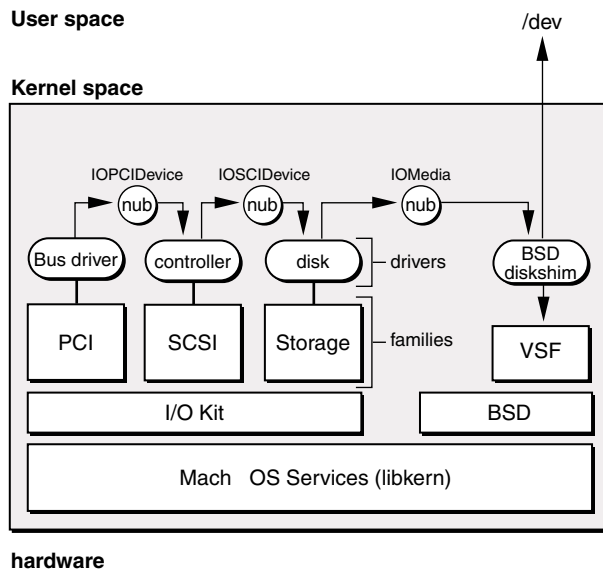
Families are implemented in C code and C++ classes. Families may include headers, libraries, sample code, test harnesses, and documentation. If it seems more familiar, however, a family can simply be thought of as a library.

A **nub** supports dynamic configuration by providing a bridge between two drivers (and, by extension, between two families). A nub is an I/O Kit object that represents a detected, controllable entity; that is, a nub represents a device or logical service. For example, a nub may represent a disk, a disk partition, a RAID, a keyboard, or any number of similar entities.

Drivers are connected via nubs, matching drivers to devices. A nub is loaded as part of the family that instantiates it.

The figure below illustrates the I/O Kit architecture, including several example drivers and their corresponding families and nubs. Many different family and driver combinations are possible; this diagram only shows one possibility. Arrows represent data flow while connection lines without arrows represent an inheritance relationship.

Figure 5-1 I/O Architecture: Families, Drivers, and Nubs



In this example, a connection is made that spans from a disk back to the PCI bus. The connection is made in several steps.

1. The PCI Bus driver finds a PCI device and announces its presence by creating a nub. The nub is defined by the PCI family.
2. The nub has the job of identifying (matching) the correct driver and requesting that the driver be loaded. At the end of this matching process, a SCSI controller driver is found and loaded. Loading the controller driver causes all required families to be loaded as well. In this case, the SCSI family is loaded; the PCI family (also required) is already present. The SCSI controller driver is given a reference to the nub.
3. The SCSI controller driver scans the bus for SCSI devices. Upon finding a device, in this case a disk, it announces the presence of the device by creating a nub. This nub is defined by the SCSI family.
4. The nub has the job of identifying (matching) the correct driver and requesting that the driver be loaded. At the end of this matching process, a disk driver is

found and loaded. Loading the disk driver causes all required families to be loaded as well. In this case, the storage family is loaded; the SCSI family (also required) is already present. The disk driver is given a reference to the nub.

The **BSD Disk shim** provides connections between the disk driver and the file system. It also provides access to the disk via the BSD `/dev` directory. See the next section for additional details.

Families

Families provide services for many different categories of devices. For example, there are protocol families (SCSI, USB, Firewire...), storage families (disk), and network families. Families may also exist to describe more abstract services, such as human interface devices (mouse, keyboard, joy stick). When devices have features in common, the software that describes those features will most likely be found in a family.

For example, all SCSI controllers have certain things they must do, such as scanning the SCSI bus. Because this functionality has been included in the SCSI family, a developer does not need to include this code in each new SCSI controller driver. The abstractions are defined by the family, allowing all SCSI controllers to share similar features.

Families are dynamically loadable; they are loaded when needed and unloaded when no longer needed. Although some common families may be pre-loaded at system startup, all families should be considered to be dynamically loadable (and, therefore, potentially unloaded!).

Each nub provides access to the device or service that it represents, and provides such services as matching, arbitration, and power management. It is most common that a driver publishes one nub for each individual device or service it controls; however, it is also possible for a driver that vends only a single device or service to act as its own nub.

When a driver is loaded, its required families are also loaded to provide necessary, common functionality. The request to load a driver causes all of its dependent requirements (and their requirements) to be loaded first. After all requirements are met, the requested driver is loaded as well.

Note that families are loaded upon demand of the device driver, not the other way around. Occasionally, a family may already be loaded when a driver demands it;

however, this should never be assumed. To ensure that all requirements are met, each device driver should list all of its requirements in its property list.

Each device driver is in a client-provider relationship, wherein every driver must know about both the family it inherits from and the family it connects to. A SCSI controller driver, for example, must be able to communicate with both the SCSI family and the PCI family (as a client of PCI and provider of SCSI). A SCSI disk driver communicates with both the SCSI and Storage families.

Scope

All drivers exist to drive hardware; all hardware is different. Even using the reusable model provided by the I/O Kit, developers will need to be aware of the “quirks” of the hardware. The code to support those quirks will still need to be unique from driver to driver.

In contrast with “traditional” I/O models, reusable code can decrease the developer’s work substantially. When porting drivers from Mac OS 9, the Mac OS X counterparts have been up to 75% smaller.

Not all drivers benefit to this degree. Some things are difficult to abstract. The I/O Kit attempts to represent, in software, the same hierarchy that exists in hardware. When that hierarchy is difficult to represent (for example, if layering violations occur), then the I/O Kit abstractions provide less help for writing drivers.

In general, all hardware support is provided by I/O Kit entities. One notable exception to this rule is imaging devices such as printers, scanners, and cameras (although these also make some use of I/O Kit functionality). Support for imaging devices is provided by the Core Graphics Framework which is not part of the kernel.

It is also possible to write a custom driver for a device that is not abstracted by any I/O Kit family. In doing so, however, the developer no longer benefits from the services a family provides. In this case, the developer is responsible for defining an abstraction for the device, and for implementing to this abstraction, as well as for writing any client libraries necessary to allow clients to make use of the device. For

most hardware devices, an I/O Kit driver that builds on an I/O Kit family is the safest and easiest approach.

In designing the I/O Kit, one goal has been to make developers' lives easier. Unfortunately, it is not possible to make all developers' lives uniformly easy. Therefore, a second goal of the I/O Kit design is to meet the needs of the majority of developers, without getting in the way of the minority that need lower level access to the hardware.

Although most developers should be able to take advantage of an I/O Kit family, there will always be some who cannot. Some driver writers will only need to override a few things; others (rarely) will be unable to find any I/O Kit abstraction that fits their device. For this reason, the source code is always available. Developers who need to do so will be able to replace functionality and modify the classes themselves.

Presenting kernel APIs in user space

Mac OS X draws a distinction between kernel and user space. In general, applications in user space cannot interface directly with kernel APIs.

In some cases, a kernel API needs to be presented in user space. For example, a disk backup application may wish to present itself to the user as if it were a disk or tape "driver", acquiring a nub and communicating with it. Or a control panel may need to interact with system software to set monitor depth or sound volume. Other examples of user applications that may need to interact with nubs in kernel space might include scanners, printers, digital cameras, and so forth.

When a user-level application must communicate with the kernel, it does so by means of a **user client**. This allows the user-level code to communicate across the user-kernel address space boundary. Some families already provide the necessary functionality. For custom drivers that do not use I/O Kit families, however, the developer will need to write additional code.

Some families, such as the SCSI family, present a traditional "split object". The user client handles negotiation, protection, authentication, and other tasks as if it were a "real" (kernel side) driver.

Other family services are never exported to user space. These services are available only inside the kernel. One such example is the PCI family. For security reasons, direct access to PCI resources from user space is forbidden.

Any code to communicate between user space and kernel space must take advantage of one or more of the following facilities available in Mac OS X:

- BSD system calls
- Mach IPC
- Mach shared memory

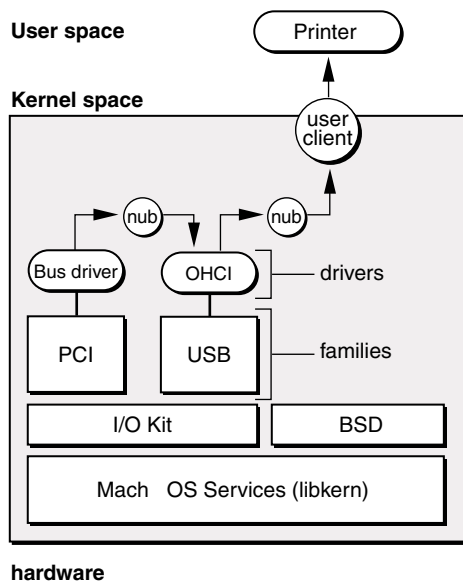
The I/O Kit mostly uses Mach IPC and Mach shared memory. In contrast, the networking and file system components of Mac OS X primarily use BSD system calls.

A user client is implemented in one or two parts. The kernel portion is usually part of an appropriate family. The (optional) user portion is linked into the application. A user client looks like a library when “viewed” from user space. From kernel space, it looks like a driver.

The BSD Disk shim is a custom driver implemented entirely in kernel code. It uses BSD system calls and the I/O Kit user client facility to export device driver interfaces into user space as BSD-style device nodes (in the `/dev` directory). The BSD disk shim also communicates with the file system and VFS stacks. Although the BSD disk shim is included in the storage family, it does not inherit directly from the storage family.

The figure below illustrates one example of a user client, in this case, a USB printer application. A user client permits the communication of raw USB commands across the user-kernel address space boundary.

Figure 5-2 User-space applications and split client drivers



6 Networking and Network Kernel Extensions

Network kernel extensions (NKEs) represent a specific case of a Mac OS X kernel extension. NKEs provide a way to extend and modify the networking infrastructure of Mac OS X dynamically, without recompiling or relinking the kernel. The effect is immediate and does not require rebooting the system.

Much of the content of this chapter has been excerpted from chapter 1 of *Inside Mac OS X: Network Kernel Extensions*. For a further (and more in-depth) reference to this topic, you should refer to that book.

NKEs allow you to

- create protocol stacks that can be loaded and unloaded dynamically and configured automatically
- create modules that can be loaded and unloaded dynamically at specific positions in the network hierarchy.

These modules can monitor network traffic, modify network traffic, and receive notification of asynchronous events at the data link and network layers from the driver layer, such as power management events and interface status changes.

The Kernel Extension Manager dynamically adds NKEs to the running Mac OS X kernel inside the kernel's address space. An installed and enabled NKE is invoked automatically, depending on its position in the sequence of protocol components, to process an incoming or outgoing packet.

As KEXTs, all NKEs provide initialization and termination routines that the Kernel Extension Manager invokes when it loads or unloads the NKE. The initialization routine handles any operations needed to complete the incorporation of the NKEs into the kernel, such as updating `protosw` and `domain` structures. Similarly, the termination routine must remove references to the NKE from these structures to

unload itself successfully. NKEs must provide a mechanism, such as a reference count, to ensure that the NKE can terminate without leaving dangling pointers.

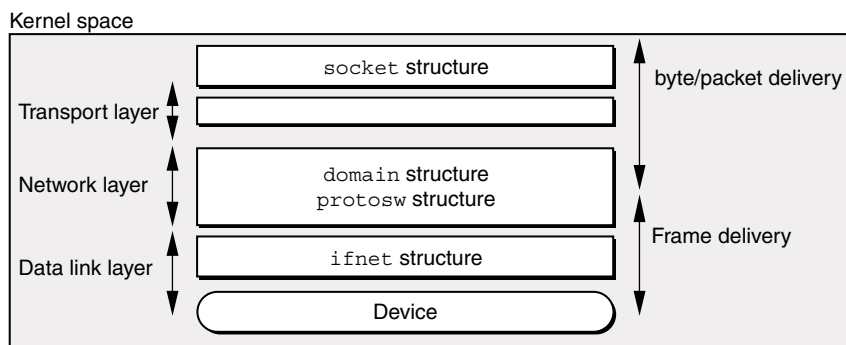
Review of 4.4BSD Network Architecture

Mac OS X is based on the 4.4BSD operating system. The following structures control the 4.4BSD network architecture:

- `socket` structure — used to keep track of network information on a per-file descriptor basis. The `socket` structure is referenced by file descriptors from user mode.
- `domain` structure — describes protocol families.
- `protosw` structure — describes protocol handlers. (A protocol handler is the implementation of a particular protocol in a protocol family.)
- `ifnet` structure — describes a network interface

None of these structures is used uniformly throughout the 4.4BSD networking infrastructure. Instead, each structure is used at a specific level, as shown below.

Figure 6-1 4.4 BSD network architecture



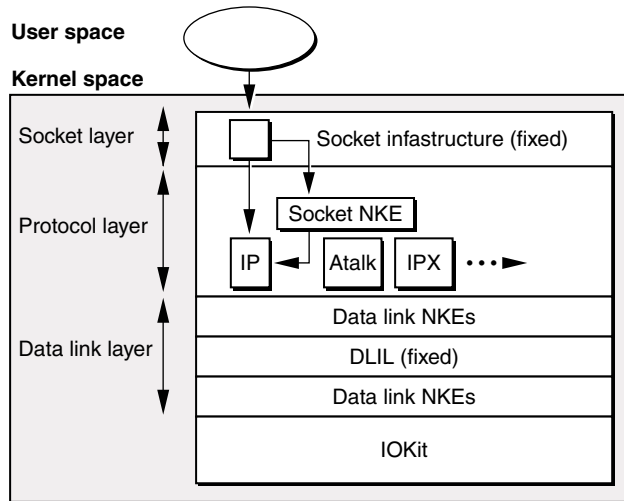
Above the network layer, packets are isolated on a per-user (per-file descriptor) basis. That is, packets are isolated based upon their ownership. Below the network layer, packets are isolated based on which device they go to (or originate from). The network layer provides a transition in how packets are viewed and processed. In the protocol stack (network layer) and the data link layer, the point of view is per-packet. Above these, in the `socket` structure, the point of view is the stream.

NKE Types

Making the 4.4BSD network architecture dynamically extensible requires several NKE types, for use at specific places in the kernel.

- `socket` NKEs — these reside between the `socket` layer and the transport protocol handlers and are invoked through a `protosw` structure. `Socket` NKEs use a new set of dispatch vectors that intercept specific `socket` and `socket` buffer utility functions.
- `protocol family` NKEs — these are collections of protocols that share a common addressing structure. Internally, a `domain` structure and a chain of `protosw` structures describe each protocol.
- `protocol handler` NKEs — these process packets for a particular protocol within the context of a protocol family. A `protosw` structure describes a protocol handler and provides the mechanism by which the handler is invoked to process incoming and outgoing packets and for invoking various control functions.
- `data link` NKEs — inserted below the protocol layer and above the network interface layer. This type of NKE can passively observe traffic as it flows in and out of the system (for example, a sniffer) or can modify the traffic (for example, encrypting or performing address translation).

The following figure summarizes the NKE architecture.

Figure 6-2 NKE architecture

NKE Modes

Socket NKEs operate in one of two modes: programmatic or global. Data link NKEs operate only in global mode.

A programmatic NKE is a socket NKE that is enabled under program control, using socket options, for a specific socket. That is, a program is responsible for enabling them on a specific socket. Programmatic NKEs must be specified by a name (a 32-bit integer handle); these should be registered with Apple. NKE handles use the same namespace as Type and Creator handles.

In contrast, global socket NKEs as well as data link NKEs are automatically enabled when they are loaded and initialized. The developer (or application) need not know the names of the global NKEs that are enabled.

Modifications to 4.4BSD Networking Architecture

To support NKEs in Mac OS X, the 4.4BSD `domain` and `protosw` structures were modified as follows:

- The `protosw` array referenced by the `domain` structure is now a linked list, thereby removing the array's upper bound. The new `max_protohdr` member defines the maximum protocol header size for the domain. The new `dom_refs` member is a reference count that is incremented when a new socket for this address family is created and is decremented when a socket for this address family is closed.
- The `protosw` structure is no longer an array. The `pr_next` member has been added to link the structures together. This change has implications for `protox` usage for `AF_INET` and `AF_ISO` input packet processing. The `pr_flags` member is an unsigned integer instead of a short. NKE hooks have been added to link NKE descriptors together.

7 File Systems and VFS Stacks

Mac OS X provides “out-of-the-box” support for several different file systems. These include Mac OS Extended Format (HFS+), the BSD standard file system format (UFS), NFS (an industry standard for networked file systems) and ISO 9660 (used for CD-ROM).

Support is also included for reading the older, Mac OS Standard Format (HFS) file-system type; however, you should not plan to format new volumes using Mac OS Standard Format. Mac OS X cannot boot from these file systems, nor does the Mac OS Standard format provide some of the information required by Mac OS X.

Mac OS X boots and “roots” from Mac OS Extended Format. That is, Mac OS X can mount a Mac OS Extended Format volume and use it as the primary or root file system. The Extended Format file system type provides many of the same characteristics as Mac OS Standard Format but adds additional support for modern features such as file permissions, longer file names, Unicode, both hard and symbolic links, and larger disk sizes.

Other file systems can be mounted, allowing users to gain access to additional volume formats and features. For example, UFS provides case sensitivity and other characteristics that may be expected by BSD commands. In contrast, Mac OS Extended Format is case-insensitive (but case-preserving).

NFS provides access to network servers as if they were locally mounted file systems. The Carbon application environment mimics many expected behaviors of Mac OS Extended Format on top of both UFS and NFS. These include such characteristics as Finder Info, file ID access, and aliases.

By using the Mac OS X Virtual File System (VFS) capability and writing kernel extensions, developers can add support for other file systems. Examples of file systems that are not currently supported in Mac OS X but that third party developers may wish to add to the system include the Andrew file system (AFS)

and the Windows NT File system (NTFS). If you want to support a new volume format or networking protocol, you'll need to write a file system kernel extension.

Working with the File System

In Mac OS X, the **vnode** structure provides the internal representation of a file or directory (folder). There is a unique vnode allocated for each active file or folder, including the root.

Within a file system, operations on specific files and directories are implemented via vnodes and **VOP** calls. VOP calls are used for operations on individual files or directories (such as open, close, read, or write). Examples include `VOP_OPEN` to open a file and `VOP_READ` to read file contents. In contrast, file system-wide operations are implemented using VFS calls. VFS calls are primarily used for operations on entire file systems; examples include `VFS_MOUNT` and `VFS_UNMOUNT` to mount or unmount a file system, respectively. File system writers need to provide stubs for each of these sets of calls.

Supporting a new volume format requires implementing a new file system type. However, it is not always necessary to implement a new file system type in order to change the way in which a user interacts with files. VFS stacks allow developers to create and layer new capabilities onto an existing file system type.

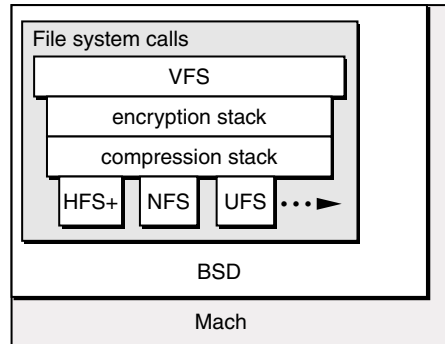
VFS stacks provide filters between the user and the underlying file system. As implied by the diagram, VFS stack can run on top of any type of file system. If your application does not need to support a volume format or networking protocol, but does need to intercept data going into or out of the file system, implementing a VFS stack may be the better choice.

For example, VFS stacks may be used in the following sorts of application areas

- virus checking — automatically check a file for viruses before reading its data
- compression — perform compression or decompression “on-the-fly” when opening (reading) and writing files
- encryption — automatically encrypt a file as it is written, then decrypt it (with a password) when it is opened or read

The following figure illustrates the file system architecture with several example VFS stacks and file systems shown.

Figure 7-1 File systems and VFS stacks



When writing a VFS stack, the developer must create a stub for each vnode operation. In some cases, the stub will simply call the routine of the same name in the underlying layer. Note that stacks may be implemented directly on top of a file system or on top of other stacks, so a developer cannot be sure exactly what the underlying layer will do with a given VFS or VOP call.

In other cases, such as when creating, reading, or writing a file, the VFS stack developer will intercept a call rather than simply passing it to the underlying layer. For example, an encryption stack may intercept read and write calls in order to add encryption or decryption filters. A virus checking stack might intercept the open and read calls.

VFS stacks are KEXTs. The Kernel Extension Manager dynamically adds VFS stacks and support for additional file-system types to the running Mac OS X kernel as part of the kernel's address space. An installed and enabled file system can be mounted automatically or manually. Further file access goes through that file system's calls.

A Politically Correct Example

The “Politically Correct File System” is an example of a VFS stack. In this example, all calls are ignored (passed to the underlying layer) except for those that create, read, or write a file (or folder).

Upon receiving a request to create a file or folder, the Politically Correct (PC) stack intercepts the call before it can be executed by the underlying file system. The PC version of the create call checks the requested file name against a table of names. If the name is deemed “politically incorrect”, for example if the user chooses to name a file “vulgarity”, the PC create call will choose a more pleasing name, for example, “politeness”. The new name will be passed to the create routine of the underlying file system.

Similarly, when a user opens a file to read or write it, such as with a text editor, the PC read and write routines first examine the data buffer, possibly substituting preferred words and phrases for their undesirable counterparts. After the substitutions are made, the buffer would be handed to the underlying routine, for example to write the data to disk.

Thus, if a user attempted to save a file containing a sentence such as

The beleaguered computer company's woes continue, despite rising stock prices.

the PC write routine might intercept and filter this sentence to a more desirable version

The aspiring computer company's joys continue, due to rising stock prices.

8 Extending the Kernel

KEXTs

As discussed in chapter 2, Kernel Architecture, Mac OS X provides a kernel extension mechanism as a means of allowing dynamic loading of code into the kernel, without the need to recompile or relink. Because these kernel extensions (KEXTs) provide both modularity and dynamic loadability, they are a natural choice for any relatively self-contained service that requires access to kernel internal interfaces.

Because KEXTs run in supervisor mode in the kernel's address space, they are also harder to write and debug than user-level modules, and must conform to strict guidelines. Further, kernel resources are wired (permanently resident in memory) and are thus more costly to use than resources in a user-space task of equivalent functionality.

In addition, although memory protection keeps applications from crashing the system, no such safeguards are in place inside the kernel. A badly behaved kernel extension in Mac OS X can cause more trouble than a badly behaved application or extension could in Mac OS 8 or 9.

Bugs in KEXTs can have far more severe consequences than bugs in user-level code. For example, a memory access error in a user application will, at worst, cause that application to crash. In contrast, a memory access error in a KEXT will cause a system **panic**, crashing the operating system.

Finally, for security reasons, some customers may not permit or will restrict use of third-party KEXTs. As a result, use of KEXTs is strongly discouraged in situations

where user-level solutions are feasible. Mac OS X guarantees that user threads are just as efficient as kernel threads, so efficiency should not be an issue. Unless your application requires low-level access to kernel interfaces or the data stream, you should use a higher level of abstraction when developing code for Mac OS X.

When you are trying to determine if a piece of code should be a KEXT, the default answer is generally *no*. In particular, if your code was a system extension in Mac OS 8 or 9, that does not imply that it must necessarily be a kernel extension in Mac OS X. There are only a few good reasons for a developer to write a kernel extension:

- Your code needs to take a primary interrupt, that is, something in the hardware needs to interrupt the CPU.
- The primary client of your code is inside the kernel, for example, a block device whose primary client is a file system.
- A sufficiently large number of running applications require a resource that your code provides, for example, you have written a file system stack.
- Your code needs to multiplex between multiple client applications that require high speed, excellent synchronization, or low latency.

If your code does not meet any of the above criteria, you should consider developing it as a library or a user-level daemon, or using one of the user-level plug-in architectures (such as QuickTime components or the Core Graphics Framework) instead of writing a kernel extension.

For developers writing device drivers or code to support a new volume format or networking protocol, however, KEXTs may be the only feasible solution. Fortunately, while KEXTs may be more difficult to write than user-space code, several tools and procedures are available to enhance the development and debugging process. See the section on [Debugging](#) for more information.

This chapter provides a conceptual overview of KEXTs and how to create them. If you are interested in building a simple KEXT, see the Apple tutorials referenced in the Bibliography. These provide step by step instructions for creating a simple, generic KEXT or a basic I/O Kit driver.

Implementation

KEXTs are implemented as **bundles**, folders that the Finder treats as single files. See [the Bundles chapter in *Inside Mac OS X: System Overview*](#) for a discussion of bundles. The KEXT bundle can contain the following:

- **Information property list** — a text file that describes the contents, settings, and requirements of the KEXT. This file is required. A KEXT bundle need contain nothing more than this file, although most KEXTs contain one or more kernel modules as well. See [the Software Configuration chapter in *Inside Mac OS X: System Overview*](#) for further information about property lists.
- **Kernel modules** — a file (or files) in Mach-O format, containing the actual binary code used by the KEXT. A kernel module (or **KMOD**) represents the minimum unit of code that can be loaded into the kernel. A KEXT can contain zero or more KMODs. If no KMODs are included, the information property list file must contain a reference to at least one module in another KEXT and change its default settings.
- **Resources** — for example, icons or localization dictionaries. Resources are optional; they may be useful for KEXTs that need to display a dialog box or menu. At present, no resources are explicitly defined for use with KEXTs.

Dependencies

Any KMOD can declare that it is dependent upon any other KMOD. The developer lists these dependencies in the “Requires” field of the module’s property list file.

Before a KMOD is loaded, all of its requirements will be checked. Those required modules (and their requirements) will be loaded first, iterating back through the lists until there are no more required modules to load. Only after all requirements are met, will the requested KMOD be loaded as well.

For example, device drivers (a type of KEXT) are dependent upon (require) certain families (another type of KEXT). When a driver is loaded, its required families are also loaded to provide necessary, common functionality. To ensure that all requirements are met, each device driver should list all of its requirements (families and other drivers) in its property list. See chapter 5, [Device Drivers and the I/O Kit](#), for an explanation of drivers and families.

It is important to list all dependencies for each KMOD. If your KEXT fails to do so, your KMOD may not load due to unrecognized symbols, thus rendering the KEXT useless. Dependencies in KMODs can be considered analogous to required header files or libraries in code development; in fact, the Kernel Extension Manager uses the standard linker to resolve KMOD requirements.

Building and Testing your KEXT

After creating the necessary property list and C (or C++) source files, you use Project Builder to build your KEXT as well. Any errors in the source code will be brought to your attention during the build and you will be given the chance to edit your source files and try again.

To test your KEXT, however, you will need to leave Project Builder and work in the Terminal application (or in **console** mode). In console mode, all system messages are written directly to your screen, as well as to a log file (`/var/log/system.log`). If you work in the Terminal application, you'll need to view system messages in the log file. You'll also need to log in to the **root** account (or use the `su` command), since only the root account can load kernel extensions.

When testing your KEXT, you can load and unload it manually, as well as check the load status. You can use the `kextload` command to load any KEXT. This command handles matching for I/O Kit drivers, then calls `kmodload`. If you are not working with the I/O Kit you can run `kmodload` directly. Manual pages for these, as well as the `kmodunload` and `kmodstat` commands, are included in Mac OS X.

Note that these commands are only useful when developing a KEXT. Eventually, after it has been tested and debugged, you'll install your KEXT in one of the standard places (see [Installed KEXTs](#) for details). Then, it will be loaded and unloaded automatically at system startup and shutdown or whenever it is needed (such as when a new device is detected).

Debugging

KEXT debugging can be complicated. Before you can debug a KEXT, you must first enable kernel debugging, as Mac OS X is not normally configured to permit debugging the kernel. Only the root account can enable kernel debugging, and you'll need to reboot Mac OS X for the changes to take effect.

Kernel debugging is performed using two Mac OS X machines, called the “development” and “target” machines. These machines must be connected over a reliable network connection on the same subnet (or within a single local network). Specifically, there must not be any intervening IP routers or other devices which could make hardware-based Ethernet addressing impossible.

The KEXT is registered (and the KMODs loaded and run) on the target machine. The debugger is launched and run on the development machine. You can also rebuild your KEXT on the development machine, after you fix any errors you find.

Debugging must be performed in this fashion because you must halt the kernel temporarily in order to use the debugger. When you halt the kernel, all other processes on that machine stop. However, a debugger running remotely can continue to run and can continue to access the kernel on the target machine.

Note that bugs in KEXTs may cause the target kernel to freeze or panic. When this happens, you may not be able to continue debugging, even over a remote connection; you will have to reboot the target and start over, setting a breakpoint just before the code where the KEXT crashed and working very carefully up to the crash point.

KEXTs are debugged using **GDB**, a source-level debugger with a command-line interface. You will need to work in the Terminal application to run GDB. For detailed information about using GDB, see the documentation included with Mac OS X. You can also use the `help` command from within GDB.

Because KEXT debugging happens at such a low level, you won't be able to take advantage of all features of GDB. For example:

- You can't call a function or method in a KEXT.
- You can't debug interrupt routines.

Use care that you do not halt the kernel for too long when you are debugging (for example, when you set breakpoints). In a short time, internal inconsistencies can

appear that will cause the target kernel to panic or freeze, forcing you to reboot the target machine.

Installed KEXTs

The Kernel Extension Manager (KEXT Manager) is responsible for loading and unloading all installed KMODs (commands such as `kextload` are only used during development). Installed KMODs are dynamically added to the running Mac OS X kernel as part of the kernel's address space. An installed and enabled KMOD is invoked as needed.

Important

Note that KEXTs are only wrappers (bundles) around a property list, KMODs (or references to KMODs), and optional resources. The KEXT describes what is to be loaded; it is the KMODs that are actually loaded.

The Kernel Extension Manager must be able to find the KEXTs (KMODs) it will access, before it can load them. KEXTs are usually installed in the Extensions folder (at `/System/Libraries/Extensions.`) The Kernel Extension Manager (in the form of a **daemon**, `kextd`), always checks here. KEXTs can also be installed in several other locations:

- in ROM
- in the Driver partition on a disk
- inside an application bundle

The last location allows an application to register KEXTs without the need to install them permanently, elsewhere within the system hierarchy. This may be more convenient and allows the KMOD to be associated with a specific, running, application. When it starts, the application can call the Kernel Extension Manager and register a KEXT.

For example, a network packet sniffer application might employ a network kernel extension (NKE). A tape backup application would require that a tape driver be loaded during the duration of the backup process. When the application exits, the kernel extension is no longer needed and can be unloaded.

Extending the Kernel

Note that, although the application is responsible for registering the KEXT, this is no guarantee that the corresponding KMODs will actually ever be loaded. It is still up to a kernel component, such as the I/O Kit, to determine a need, such as matching a piece of hardware to a desired driver, and tell the KEXT Manager to load the appropriate KMODs (and their dependencies).

Extending the Kernel

9 Kernel Services

The Kernel as Service Provider

In a typical preemptive multitasking operating system such as Mac OS X, FreeBSD, or Linux, user applications are not allowed direct access to shared resources such as RAM, disks, printers, and other devices. Instead, the kernel provides controlled access to these resources, and can thus be viewed as a service provider.

Recall that each application exists in its own (user) address space and that the kernel exists in a separate (kernel) address space. Privileged operations, such as opening a file, initiating network traffic, or shutting down the computer, are performed in kernel space and are thus possible only to the kernel.

Applications that need to have privileged operations performed must request the appropriate services from the kernel. The kernel provides these operations as services to the processes, mapping any associated parameters in and out of user space.

Application processes include applications that are explicitly launched and run by the user, as well as various system processes, such as **daemons**, that keep the system running smoothly.

Any code to communicate between user space and kernel space must take advantage of one or more of the following facilities available in Mac OS X:

- BSD system calls
- Mach IPC

■ Shared memory

In Mac OS X, where the kernel itself is modular, interaction between the various kernel components is also in the form of services. Each component, such as Mach, networking, or the file system, is therefore both a provider of services to applications and other components as well as a client of kernel services itself. Kernel space, however, is a single address space; memory is shared between kernel components. Thus, kernel components are able to communicate more freely with each other than with applications in user space.

Available Services

Many of the most well-used kernel services are described below. For each service, the provider component is named as well as the client component(s). A brief description is also given. For more complete information, see the available documentation for the component itself.

In the API listings below, header files are listed as they would be included in real code. The default compiler flags should locate the correct file in the “well known places”.

In addition, the following header files are assumed to be included at all times:

```
#include <sys/param.h>      /* useful defines and limits */
#include <sys/types.h>       /* exported data types */
#include <sys/systm.h>       /* "systm" and NOT "system". prototypes */
#include <libkern/libkern.h> /* more prototypes */
```

BSD Disk Shim

Provider: I/O Kit

Clients: BSD, File systems

The BSD Disk shim uses BSD system calls and the I/O Kit user client facility to export device driver interfaces into user space as BSD-style device nodes in the

`/dev` directory. The BSD disk shim also communicates with the file system and VFS stacks. Support for user processes is provided via `devfs`.

Events

Provider: Mach

Clients: All kernel components, user processes

Specific services include port notifications, notification ports, and notification events.

Exceptions, Traps

Provider: Mach

Clients: BSD, user processes

This service support BSD signals, interrupts, and debugging, as well as various system calls that can be accessed by user processes.

Families

Provider: I/O Kit

Clients: BSD, File systems, Networking, user processes

APIs:

`IOKit/**/*`

This service provides APIs for I/O Kit families, including support for networking, block, graphics, firewire, USB, human interface, and many other device categories.

File Descriptor Management

Provider: BSD

Clients: File systems, Networking, user processes

Kernel Services

APIs:

```
#include <sys/filedesc.h>
#include <sys/file.h>
```

File descriptors provide per-process unique, non-negative integers that are used to identify an open file (or socket). For user processes, all interaction with files is done via file descriptors. File descriptors are also used for access and manipulation of POSIX semaphores and POSIX shared memory.

Device Driver Management

Provider: I/O Kit

Clients: I/O Kit KEXTs (Families and Drivers), user processes

APIs:

```
IOKit/**/*
```

These services support device driver instantiation, matching, service notification. Family APIs publish services; drivers use devices.

Host Manipulation and Inquiry

Provider: Mach

Clients: All kernel components

These services are used to get and set host-based information, such as page size and processor count.

Inter-process Communication (IPC)

Provider: Mach

Clients: All kernel components

This service provides various specialized forms of communication between tasks (processes) on the local machine. The particular form of IPC in use dictates how

(and if) data is processed. Specific services include: send and receive operations as well as primitives for servicing ports and/or port sets. See also: Port Right Management, Tasks and Threads Management, Virtual Memory, and Low-Level Synchronization Services

Kernel Loadable Module Support

Provider: I/O Kit

Clients: File systems, Networking, Loadable modules

Tools: See the man pages for the following utility programs:

```
man kextload
man kmodload
man kmodstat
man kmodunload
```

This service provides support for loading and unloading KEXTs.

Kernel Tracing

Provider: BSD

Clients: Mach, I/O Kit, File systems, Networking, Loadable modules

APIs:

```
#include <sys/kdebug.h>
```

This service provides information for performance analysis and debugging support as well as trace points for user processes.

Lock Management

Provider: BSD

Clients: File systems, Networking, loadable modules, user processes

APIs:

Kernel Services

```
#include <sys/lock.h>
```

BSD, files system and networking code should use this service for management of locking operations. Note that this API is quite different from the one defined in `osfmk/kern/lock.h`.

Mach Interface Generator (MIG)

Provider: Mach

Clients: All kernel components, user processes (Project Builder)

MIG is used to specify IPC formats that are valid on a given port. It is used mostly in Remote-Procedure Call (RPC) situations, but supports other forms of communication as well. MIG also provides a set of runtime services for dispatching incoming communications to the appropriate handler. Project Builder has special rules and targets for generating stubs for both sides of the MIG interface.

mbuf Management

Provider: Networking

Clients: NKEs (third-party), File systems

APIs:

```
#include <sys/mbuf.h>
```

These services provide support for the `mbuf` data structure which is used to manage I/O for network devices.

Memory and Address Space Management

Provider: Mach

Clients: All kernel components, user processes

Specific services include virtual memory management, address space allocation, page read and write, external memory managers (EMMI), and memory objects.

Port Right Management

Provider: Mach

Clients: All kernel components

Port right ownership is the fundamental security mechanism within Mach. Specific services include creation and destruction, reference management, copying, explicit insertion and removal from other tasks, passing via IPC, grouping rights into sets, and requesting asynchronous notifications about changes in a port's status. See also: Task and Thread Management, IPC.

Processor Management

Provider: Mach

Clients: All kernel components

These services provide low level hardware support, including processor start, processor stop, and power management.

Registry

Provider: I/O Kit

Clients: I/O Kit Family APIs, user processes.

APIs:

```
#include <IOKit/IORegistryEntry.h>
```

These services support I/O Kit families in publishing devices or services and device information.

Queue Management

Provider: BSD

Clients: File systems, Networking, loadable modules, user processes

APIs:

```
#include <sys/queue.h>
```

BSD, file systems and networking code use this service for queue management. It provides support for singly and doubly linked lists and queues. Note that there are subtle differences between this API and the queues found in `osfmk/kern/queue.h`

Socket management

Provider: Networking

Clients: NKEs (third-party), File systems, user processes

APIs:

```
#include <sys/socket.h>
#include <sys/socketvar.h>
```

These services provide support for the management of sockets.

Network kernel extension support

Provider: Networking

Clients: NKEs (third-party)

APIs:

```
#include <net/kext_net.h>
```

These services provide general support for network kernel extensions.

Scheduling

Provider: Mach

Clients: BSD

Specific services include priority-based thread scheduling, preemption, and processor resource allocation, based on the following policies: timesharing, round-robin, and **FIFO** fixed priority.

Synchronization primitives (low level)

Provider: Mach

Clients: All kernel components, IPC services for exporting to user space.

This Mach service provides low-level implementation support for basic asynchronous primitives (wait queues, semaphores) as well as basic locking primitives (machine-specific locks, spin locks, mutexes, shared/exclusive read/write locks).

Synchronization primitives

Provider: BSD

Clients: File systems, Networking, loadable modules, user processes

APIs:

```
#include <sys/proc.h>
#include <machine/spl.h>
```

This BSD service provides higher level support for `sleep()` and `wakeup()` calls as well as SPLs.

sysctl

Provider: BSD

Clients: File systems, Networking, loadable modules

Kernel Services

APIs:

```
#include <sys/sysctl.h>
```

This service provides a formalized interface for kernel global manipulation and tuning.

Task and thread management

Provider: Mach

Clients: BSD

This service provides the underlying implementation for BSD process management; a process is based on one Mach task and one or more Mach threads. A task is the unit of resource ownership. A thread is an independently schedulable execution path.

Timing Services

Provider: Mach

Clients: BSD, user processes

The kernel provides several different timing services to user processes. Timing services support profiling, statistics gathering, and various types of timers, as well as current date and time-of-day functionality.

VFS Infrastructure

Provider: BSD

Clients: File system

APIs:

```
#include <sys/buf.h>
#include <sys/namei.h>
#include <sys/vnode.h>
#include <sys/vnode_if.h>
```

Kernel Services

```
#include <vfs/vfs_support.h>
```

This service provides VFS management routines and default library routines in support of virtual file system functionality.

Vnode management

Provider: BSD

Clients: File system

APIs:

```
#include <sys/vnode.h>
```

This services provides allocation, referencing, and serialization functionality in support of vnode management.

Zone allocator

Provider: Mach

Clients: BSD, Networking

This services provides support for efficient kernel memory allocation.

Kernel Services

10 Glossary

address space

The address space of a process describes the range of memory (both physical and virtual) that it uses while running. In Mac OS X, processes do not share address space.

Apple Public Source License

Apple's Open Source license, available at <http://www.publicsource.apple.com>. The Darwin OS is distributed under this license. See also Open Source.

AppleTalk

A suite of network protocols that is standard on Macintosh computers and can be integrated with other network systems, such as the Internet.

ASCII

American Standard Code for Information Interchange. A 7-bit character set (commonly represented using 8-bits) that defines 128 unique character codes. See also Unicode.

bootstrap port

A Mach port to which a new task can send a message that will return any other system service ports the task needs.

BSD

Berkeley Software Distribution. Formerly known as the Berkeley version of UNIX, BSD is now simply called the BSD operating system. The BSD portion of Mac OS X is based on 4.4BSD Lite 2 and FreeBSD, a "flavor" of 4.4BSD.

BSD disk shim

Part of the I/O Kit Storage Family. The BSD Disk Shim provides access to all storage device being managed by I/O Kit drivers via traditional BSD device nodes.

Glossary

bundle

A directory in the file system that stores executable code and the software resources related to that code. Applications, plug-ins, and frameworks are types of bundles. Except for frameworks, bundles are file packages, presented by the Finder as a single file.

Carbon

An application environment on Mac OS X that features a set of programming interfaces derived from earlier versions of the Mac OS. The Carbon APIs have been modified to work properly with Mac OS X, especially with the foundation of the operating system, the kernel environment. Carbon applications can run on Mac OS X, Mac OS 9, and all versions of Mac OS 8 later than Mac OS 8.1.

CFM

Code Fragment Manager.

Classic

An application environment on Mac OS X that lets you run non-Carbon legacy Mac OS software. It supports programs built for both Power PC and 68k processor architectures and is fully integrated with the Finder and the other application environments.

Cocoa

An advanced object-oriented development platform on Mac OS X. Cocoa is a set of frameworks with programming interfaces in both Java and Objective-C. It is based on the integration of OPENSTEP, Apple technologies, and Java.

client-provider relationship

In the I/O Kit, the relationship between a driver and the family it inherits from (client relationship) and the family it connects to (provider relationship).

condition variable

A type of variable provided by the POSIX threads functions to help synchronize the threads in a task.

console

A special window that displays system log messages, as well as output written to the standard error and standard output streams by applications launched from the Finder. Also, an application by the same name, that displays this information.

Glossary

cooperative multitasking

A multitasking environment in which a running program can receive processing time only if other programs allow it; each application must give up control of the processor “cooperatively” in order to allow others to run. Mac OS 8 and 9 are cooperative multitasking environments. See also preemptive multitasking.

Data Link Interface Layer (DLIL)

The part of the Darwin/Mac OS X kernel’s networking infrastructure that provides the interface between protocol handling and network device drivers in I/O Kit. A generalization of the BSD “ifnet” architecture.

Darwin

Another name for the Mac OS X core operating system. The Darwin kernel is synonymous with the Mac OS X kernel.

Darwin OS

Darwin OS is the name of an Open Source product that includes the Darwin kernel, the BSD commands and C libraries, and several additional features.

daemon

A long-lived process, usually without a visible user interface, that performs a system-related service. Daemons are usually spawned automatically by the system, and may either live forever or be regenerated at intervals.

demand paging

An operating system facility that brings pages of data from disk into physical memory only as they are needed.

device driver

A component of an operating system that deals with getting data to and from a device, as well as control of that device.

DMA

Direct memory access; a means of transferring data between host memory and a peripheral device without involving the host processor.

DVD

Digital Versatile Disc or Digital Video Disc. An optical storage medium that provides greater capacity and bandwidth than CD-ROM; DVDs are frequently used for multimedia as well as data storage.

dyld

Dynamic link editor.

Glossary

Embedded C++

A dialect of C++ that supports a subset of all C++ features. Many of the more “exotic” features of C++ are not supported; these include such features as multiple inheritance and runtime type checking.

EMMI

Mach’s External Memory Management Interface. See also external pager.

Ethernet

A high-speed local area network technology.

exception

An interruption to the normal flow of program control caused by the program itself.

exception port

A Mach port on which a task or thread receives messages when exceptions occur.

external pager

A module that manages the relationship between virtual memory and its backing store. External pagers are clients of Mach’s EMMI. They may be either in the kernel or in user space. The built-in pagers in Mac OS X are the default pager and the vnode pager. See also EMMI.

family

In the I/O Kit, a family defines a collection of software abstractions that are common to all devices (instances) of a particular category (device class). Families provide functionality and services to drivers. Drivers provide the implementation behind the abstraction

FAT

An acronym for File Allocation Table, a data structure used in the MS-DOS file system. Also synonymous with the file system that uses it. The FAT file system is also used as part of Microsoft Windows and has been adopted for use inside devices such as digital cameras. In lower case, “fat files” are Mach-O files containing object code for more than one machine architecture.

FIFO

First In First Out. A data processing scheme in which data is read in the order in which it was written, processes are run in the order in which they were scheduled, and so forth.

Glossary

file descriptor

A per-process unique, non-negative integer used to identify an open file (or socket).

firewall

Software (or a computer running such software) that prevents unauthorized access to a network by users outside of the network.

fork

A stream of data that can be opened and accessed individually under a common file name. The Macintosh Standard and Extended file systems store a separate “data” fork and a “resource” fork as part of every file; data in each fork can be accessed and manipulated independently of the other. In BSD, `fork` is a system call that creates a new process.

framework

A basic structure that holds the parts of some thing together. In Mac OS X, specifically, a bundle containing a dynamic shared library and associated resources, including image files, header files, and documentation.

FreeBSD

A “flavor” of the 4.4BSD operating system.

GDB

GNU Debugger. GDB is a powerful, source-level debugger with a command line interface. GDB is a popular Open Source debugger and is included with Mac OS X.

host

The computer that’s running (is host to) a particular program. The term is usually used to refer to a computer on a network.

host processor

The microprocessor on which an application program resides. When an application is running, the host processor may call other, peripheral microprocessors, such as a digital signal processor, to perform specialized operations.

HFS

Hierarchical File System. The Mac OS Standard file system format, used to represent a collection of files as a hierarchy of directories (folders), each of which may contain either files or folders themselves.

HFS+

Hierarchical File System Plus. The Mac OS Extended file system format. This file system format was introduced as part of Mac OS 8.1,

Glossary

adding support for file names longer than 31 characters, Unicode representation of file and directory names, and efficient operation on very large disks.

IDE

Interactive Development Environment; alternatively, Integrated Development Environment. An application or set of tools that allows a programmer to write, compile, edit, and perhaps test and debug within an integrated, interactive environment.

inheritance attribute

In Mach, a value indicating the degree to which a parent process and its child process share the parent process's address space. A memory page can be inherited copy-on-write, shared, or not at all.

in-line data

Data that's included directly in a Mach message, as opposed to referred to by a pointer. See also out-of-line data.

I/O

Input/output; the sending and retrieving of information into the memory of a program, usually to and from a file or a peripheral device.

Info Plist

See information property list.

information property list

A special form of property list with predefined keys for specifying basic bundle attributes and information of interest, such as supported document types and offered services. See also bundle, property list.

IPC

Inter-process communication; the transfer of information between processes.

KDK

Kernel Development Kit.

kernel

The complete Mac OS X core operating system environment that includes Mach, BSD, the I/O Kit, File systems, and Networking components.

kernel extension

See KEXT.

Glossary

kernel port

A Mach port whose receive right is held by the kernel. See also task port; thread port.

KEXT

Kernel EXTension. Kernel extensions extend the functionality of the kernel. The I/O Kit, File system, and Networking components are designed to allow and expect the creation and use of KEXTs.

KMOD

Kernel module. A KMOD is the minimum unit of code that can be loaded into the kernel. See also KEXT.

Mach

The lowest level of the Mac OS X kernel. Mach provides such basic services and abstractions as threads, tasks, ports, IPC, scheduling, physical and virtual address space management, VM, and timers.

Mach-O

Mach Object file format.

Mach factor

A measurement of how busy a Mach-based system (such as Mac OS X) is. Unlike a load average (as used in Linux or BSD systems), higher Mach factors mean the system is less busy.

Mach server

A task that provides services to clients, using a MiG-generated RPC interface.

main thread

By default, a process has one thread, the main thread. If a process has multiple threads, the main thread is the first thread in the process. A user process can use the POSIX thread API to create other user threads.

makefile

A makefile details the files, dependencies, and rules by which an executable application is built or by which a set of programs may be run.

memory protection

A system of memory management in which programs are prevented from being able to modify or corrupt the memory partition of another program. Mac OS 8 and 9 do not have memory protection; Mac OS X does.

Glossary

memory-mapped files

A facility that maps virtual memory onto a physical file. Thereafter, any access to that part of virtual memory causes the corresponding page of the physical file to be accessed. The contents of the file can be changed by changing the contents in memory.

message

In Mach, a message consists of a header and a variable-length body; some operating system services are invoked by passing a message from a thread to the Mach port representing the task that provides the desired service.

microkernel

A kernel implementing a minimal set of abstractions. Typically, higher-level OS services such as file systems and device drivers are implemented in layers above a microkernel, possibly in trusted user-mode servers. See also monolithic kernel.

MiG

Mach's message interface generator. MiG provides a procedure call interface to Mach's system of interprocess messaging.

monolithic kernel

A kernel architecture in which all pieces of the kernel are closely intertwined. A monolithic kernel provides substantial performance improvements; however, it is difficult to evolve the individual components independently. The Darwin Core OS is not a monolithic kernel. Rather, it is a hybrid of the monolithic and microkernel models.

multicast

A process in which a single packet may be addressed to multiple recipients. Multicast is used, for example, in streaming video, in which many megabytes of data are sent over the network.

multi-homing

The ability to have multiple network addresses in one computer. For example, multi-homing might be used to create a system in which one address is used to talk to hosts outside a firewall and the other to talk to hosts inside; the computer provides facilities for passing information between the two.

multitasking

Describes an operating system that allows the concurrent execution of multiple programs. Mac OS X uses preemptive multitasking. Mac OS 8 and 9 use cooperative multitasking.

Glossary

mutex variable

Mutual exclusion variable; a type of variable provided by the POSIX threads functions to help protect critical regions in a multiple-thread task.

NAT

Network Address Translation. A scheme that transforms network packets at a gateway so network addresses that are valid on one side of the gateway are translated into addresses that are valid on the other side.

network

A group of hosts that can directly communicate with each other.

NFS

Network File System. An NFS file server allows users on the network to share files as if they were on their own local disk.

NKE

Network Kernel Extension. NKEs provide a way to extend and modify the networking infrastructure of Mac OS X dynamically, without recompiling or relinking the kernel. The effect is immediate and does not require rebooting the system.

NMI

Non-maskable interrupt; an interrupt produced by a particular keyboard sequence or button. It can be used to interrupt a hung system.

notify port

A Mach port on which a task receives messages from the kernel advising it of changes in port access rights and of the status of messages it has sent.

nonsimple message

In Mach, a message that contains either a reference to a port or a pointer to data.

nub

An I/O Kit object that represents a device or logical service. Each nub provides access to the device or service it represents, and provides such services as matching, arbitration, and power management. It is most common that a driver publishes one nub for each individual device or service it controls; it is possible for a driver that vends only a single device or service to act as its own nub.

Glossary

NVRAM

Non-volatile RAM. RAM storage that retains its state even when the power is off.

Open Transport

A communications architecture for implementing network protocols and other communication features on computers running the Mac OS. Open Transport provides a set of programming interfaces that supports, among other things, both the AppleTalk and TCP/IP protocols.

Open Source

A definition of software that includes freely available access to source code, redistribution, modification, and derived works. The full definition is available at www.opensource.org.

out-of-line data

Data that's passed by reference in a Mach message, as opposed to being included in the message. See also in-line data.

packet

An individual piece of information sent on a network.

physical address

An address to which a hardware device, such as a memory chip, can directly respond. Programs, including the Mach kernel, use virtual addresses that are translated to physical addresses by mapping hardware controlled by the Mach kernel.

PEF

Preferred Executable Format for MacOS.

POSIX

The Portable Operating System Interface. An operating system interface standardization effort supported by ISO/IEC, IEEE and The Open Group.

port

In Mach, a secure uni-directional channel for communication between tasks running on a single system. In IP transport protocols, an integer identifier used to select a receiver for an incoming packet, or to specify the sender of an outgoing packet.

port access rights

In Mach, the ability to send to or receive from a Mach port.

Glossary

port set

In Mach, a set of zero or more Mach ports. A thread can receive messages sent to any of the ports contained in a port set by specifying the port set as a parameter to `msg_receive()`.

preemptive multitasking

A type of multitasking in which the operating system can interrupt a currently running task in order to run another task, as needed. See also cooperative multitasking.

preemption

The act of interrupting a currently running program in order to give time to another task.

priority

In Mach scheduling, a number between 0 and 127 that indicates how likely a thread is to run. The higher the thread's priority, the more likely the thread is to run. See also scheduling policy.

process

A BSD abstraction for a running program. A process' resources include a virtual address space, threads and file descriptors. In Mac OS X, a process is based on one Mach task and one or more Mach threads.

process identifier, or process ID

A number that uniquely identifies a process.

protected memory

See memory protection.

protocol handler

A network module that extracts data from input packets (giving the data to interested programs) and inserts data into output packets (giving the output packet to the appropriate network device driver).

programmed I/O

I/O in which the CPU accomplishes data transfer with explicit load and store instructions to device registers, as opposed to DMA. Byte-by-byte or word-by-word data transfer to a device. Also known as "direct I/O." See also DMA.

property list

A textual way to represent data. Elements of the property list represent data of certain types, such as arrays, dictionaries, and strings. System routines allow programs to read property lists into memory and convert the textual data representation into "real" data. See also information property list.

Glossary

Pthreads	POSIX threads implementation.
quantum	The fixed amount of time a thread or process can run before being preempted.
RAM	Random-access memory; memory that a microprocessor can either read or write to.
real time	A concept of time when using a computer. Predictable, time-critical behavior. If the user defines or initiates an event and the event occurs instantaneously, the computer is said to be operating in real time. Real-time support is important for applications such as multi-media.
receive rights	In Mach, the ability to receive messages on a Mach port. Only one task at a time can have receive rights for any one port. See also send rights.
reply port	A Mach port associated with a thread that's used in Mach remote procedure calls.
ROM	Read-only memory, that is, memory that cannot be written to.
root	Also called "superuser". A user account with special privileges. For example, only the root account can load kernel extensions.
RPC	Remote Procedure Call. In Mach, RPCs are implemented using MiG-generated messages.
SCSI	Small Computer Systems Interface. A standard connector and communications protocol used for connecting devices such as disk drives to computers.
scheduling	The determination of when each process or task runs, including assignment of start times.

Glossary

scheduling policy

In Mach, a thread's scheduling policy determines how the thread's priority is set and under what circumstances the thread runs. See also [priority](#).

send rights

In Mach, the ability to send messages to a Mach port. Many tasks can have send rights for the same port. See also [receive rights](#).

simple message

In Mach, a message that contains neither references to ports nor pointers to data.

SMP

Symmetric Multi-processing. An operating system in which two or more processors are managed by one kernel, sharing the same memory, having equal access to I/O devices, and in which any task, including kernel tasks, can run on any processor.

SPL

Set Priority Level. A request that sets the current processor priority level, the level used by the kernel to control interrupt delivery to the CPU.

socket

In BSD-derived systems, a socket refers to different entities in user and kernel operation. For a user process, a socket is a file descriptor that has been allocated using `socket(2)`. For the kernel, a socket is the data structure allocated when the kernel's implementation of the `socket(2)` call is made. In AppleTalk protocols, a socket serves the same purpose as a "port" in IP transport protocols.

stackable file systems

A file system layer that has as its input the standard VFS file system interfaces and that may call other file system layers beneath it to implement file system operations. All stackable file systems support the same interface and can be layered on top of one another to add unique functionality.

task

A Mach abstraction, consisting of a virtual address space and a port name space. A task itself performs no computation; rather, it is the framework in which threads run. See also [threads](#).

Glossary

task port

A kernel port that represents a task and is used to manipulate that task. See also kernel port; thread port.

thread

In Mach, the unit of CPU utilization. A thread consists of a program counter, a set of registers, and a stack pointer. See also task.

thread port

A kernel port that represents a thread and is used to manipulate that thread. See also kernel port; task port.

thread-safe

Code that can be executed safely by several threads simultaneously.

TCP/IP

Transmission Control Protocol/Internet Protocol. An industry standard protocol used to deliver messages between computers over the network. TCP/IP support is included in Mac OS X.

UFS

UNIX File system. An industry standard file system format used in UNIX similar operating systems such as BSD. UFS in Mac OS X is a derivative 4.4BSD UFS. Specifically, its disk layout is not compatible with other BSD UFS implementations.

UDF

Universal Disk Format. The file system format used in DVD disks.

Unicode

A 16-bit character set that defines unique character codes to characters in a wide range of languages. Unlike ASCII, which defines 128 distinct characters typically represented in 8 bits, there are as many as 65,536 distinct Unicode characters that represent the unique characters used in many foreign languages.

user client

In I/O Kit, a means of allowing user-level code to communicate across the user-kernel address space boundary, as, for example, in a printer or scanner application.

UTF-8

A Unicode Transformation Format used to represent a sequence of 16-bit Unicode characters with an equivalent sequence of 8-bit characters, none of which are zero. This sequence of characters can be represented using an ordinary C language string.

Glossary

virtual address

An address that is usable by software. Each task has its own range of virtual addresses, beginning at address zero. The Mach operating system makes the CPU hardware map these addresses onto physical memory only when necessary, using disk memory at other times. See also physical address.

VFS

Virtual File System. A set of standard internal file system interfaces and utilities that facilitate support for additional file systems. VFS provides an infrastructure for file systems built in the kernel.

VM

Virtual Memory. The use of a disk partition or a file on disk to provide the same facilities usually provided by RAM. The VM provides 32 bit (minimum) protected address space for each task and facilitates efficient sharing of that address space.

XML

Extensible Markup Language. A subset of HTML, XML provides a “metalanguage” containing rules for constructing specialized markup languages. XML users can create their own tags, making XML very flexible.

Glossary

Bibliography

Apple Mac OS X Publications

Hello Kernel: Creating a Kernel Extension With Project Builder (tutorial)

Hello IOKit: Creating a Device Driver With Project Builder (tutorial)

Inside Mac OS X: Network Kernel Extensions

Inside Mac OS X: System Architecture

General UNIX and Open Source Resources

A Quarter Century of UNIX

Peter H. Salus

Addison-Wesley, 1994, ISBN 0-201-54777-5

Berkeley Software Distribution (set)

CSRG, UC Berkeley

USENIX and O'Reilly, 1994, ISBN 1-56592-082-1

The Cathedral & the Bazaar

Musings on Linux and Open Source by an Accidental Revolutionary

Eric S. Raymond

O'Reilly & Associates, 1999, ISBN 1-56592-724-9

The New Hacker's Dictionary, 3rd. Ed.

Eric S. Raymond

MIT Press, 1996, ISBN 0-262-68092-0

Open Sources

Voices from the Open Source Revolution

Edited by Chris DiBona, Sam Ockman & Mark Stone

O'Reilly & Associates, 1999, ISBN 1-56592-582-3

Proceedings of the First Conference on Freely Redistributable
Software

Free Software Foundation

FSF, 1996, ISBN 1-882114-47-7

The UNIX Companion
Harley Hahn
Osborne/MGH, 1995, ISBN 0-07-882149-5

The UNIX Desk Reference:
The hu.man Pages
Peter Dyson
Sybex, 1996, ISBN 0-7821-1658-2

The UNIX Programming Environment
Brian W. Kernighan, Rob Pike
Prentice Hall, 1984, ISBN 0-13-937681-X
Osborne, 1996, ISBN 0-07-882189-4

Internals

Advanced Topics in UNIX:
Processes, Files, & Systems
Ronald J. Leach
Wiley, 1996, ISBN 1-57176-159-4

The Complete FreeBSD
Greg Lehey
Walnut Creek CDROM Books, 1999, ISBN 1-57176-246-9

The Design and Implementation of the 4.4BSD UNIX Operating System
Marshall Kirk McKusick, et al
Addison-Wesley, 1996, ISBN 0-201-54979-4

The Design of the UNIX Operating System
Maurice J. Bach
Prentice Hall, 1986, ISBN 0-13-201799-7

Linux Kernel Internals
Michael Beck, et al
Addison-Wesley, 1996, ISBN 0-201-87741-4

Lions' Commentary on UNIX 6th Edition with Source Code
John Lions
Peer-to-Peer, 1996, ISBN 1-57398-013-7

Panic!:
UNIX System Crash Dump Analysis
Chris Drake, Kimberly Brown
Prentice Hall, 1995, ISBN 0-13-149386-8

UNIX Internals:
The New Frontiers
Uresh Vahalia
Prentice-Hall, 1995, ISBN 0-13-101908-2

UNIX Systems for Modern Architectures:
Symmetric Multiprocessing and Caching for Kernel Programmers
Curt Schimmel
Addison-Wesley, 1994, ISBN 0-201-63338-8

Optimizing PowerPC code
Gary Kacmarcik, 1995
Addison-Wesley Publishing Company
ISBN: 0-201-40839-2

papers

Berkeley Software Architecture Manual
4.4BSD Edition
William Joy, Robert Fabry, Samuel Leffler, M. Kirk McKusick, Michael Karels
Computer Systems Research Group, Computer Science Division
Department of Electrical Engineering and Computer Science
University of California, Berkely

Mach

CMU Computer Science:
A 25th Anniversary Commemorative
Richard F. Rashid, Ed.
ACM Press, 1991, ISBN 0-201-52899-1

Load Distribution:
the Implementation of the Mach Microkernel
Dejan S. Milojevic
Vieweg, 1994, ISBN 3-528-05424-7

Programming under Mach
Boykin, et al
Addison-Wesley, 1993, ISBN 0-201-52739-1

Mach Workshop and Symposium Proceedings
USENIX

Mach Workshop Proceedings
USENIX, October 1990, no ISBN

Mach Symposium Proceedings
USENIX, November 1991, no ISBN

Mach III Symposium Proceedings
USENIX, April 1993, ISBN 1-880446-49-9

Mach 3 Documentation Series
Open Group Research Institute (RI)

Final Draft Specifications OSF/1 1.3 Engineering Release
RI, May 1993

OSF Mach Final Draft Kernel Principles
RI, May 1993

OSF Mach Final Draft Kernel Interfaces
RI, May 1993

OSF Mach Final Draft Server Writer's Guide
RI, May 1993

OSF Mach Final Draft Server Library Interfaces
RI, May 1993

Research Institute Microkernel Series
Open Group Research Institute (RI)

Operating Systems Collected Papers, Volume I
RI, March 1993

Operating Systems Collected Papers, Volume II
RI, October 1993

Operating Systems Collected Papers, Volume III
RI, April 1994

Operating Systems Collected Papers, Volume IV
RI, October 1995

USENIX Conference Papers (available online at www.usenix.org)

Mach: A New Kernel Foundation for UNIX Development
Proceedings of the Summer 1986 USENIX Conference,
Atlanta, GA.

UNIX as an Application Program
Proceedings of the Summer 1990 USENIX Conference,
Anaheim, CA.

OSF RI papers

Spec '93

OSF Mach Final Draft Kernel Interfaces

OSF Mach Final Draft Kernel Principles

OSF Mach Final Draft Server Library Interfaces

OSF Mach Final Draft Server Writer's Guide

OSF Mach Kernel Interface Changes

Spec'94

OSF RI 1994 Mach Kernel Interfaces Draft
OSF RI 1994 Mach Kernel Interfaces Draft (Part A)

OSF RI 1994 Mach Kernel Interfaces Draft (Part B)

OSF RI 1994 Mach Kernel Interfaces Draft (Part C)

Additional Papers

Debugging an object oriented system using the Mach interface

Unix File Access and Caching in a Multicomputer Environment

Untyped MIG: The Protocol

Untyped MIG: What Has Changed and Migration Guide

Towards a World-Wide Civilization of Objects

A Preemptible Mach Kernel

A Trusted, Scalable, Real-Time Operating System Environment

Mach Scheduling Framework

Networking

UNIX Network Programming,
Volume 1, Second Edition:
Networking APIs: Sockets and XTI
W. Richard Stevens
Prentice Hall, 1998.

UNIX Network Programming,
Volume 2, Second Edition:
Interprocess Communications
W. Richard Stevens
Prentice Hall, 1999.

TCP/IP Illustrated, Volume 1
The Protocols
W. Richard Stevens
Addison-Wesley, 1994

TCP/IP Illustrated, Volume 2
The Implementation
W. Richard Stevens
Addison-Wesley, 1995

TCP/IP Illustrated, Volume 3,
TCP for Transactions, HTTP, NNTP, and the UNIX Domain Protocols
W. Richard Stevens
Addison-Wesley, 1996

Operating Systems

Advanced Computer Architecture:
Parallelism, Scalability, Programmability
Kai Hwang
McGraw-Hill, 1993, ISBN 0-07-031622-8

Concurrent Systems:
An Integrated Approach to Operating Systems, Database, and
Distributed Systems
Jean Bacon
Addison-Wesley, 1993, ISBN 0-201-41677-8

Distributed Operating Systems
Andrew S. Tanenbaum
Prentice Hall, 1995, ISBN 0-13-219908-4

Distributed Operating Systems:
The Logical Design
A. Goscinski
Addison-Wesley, 1991, ISBN 0-201-41704-9

Distributed Systems, Concepts, and Designs
G. Coulouris, et al
Addison-Wesley, 1994, ISBN 0-201-62433-8

Operating System Concepts, 4th Ed.
Abraham Silberschatz, Peter Galvin
Addison-Wesley, 1994, ISBN 0-201-50480-4

POSIX

Information Technology-Portable Operating System Interface (POSIX):
System Application Program Interface (API) (C Language)
ANSI/IEEE Std 1003.1, 1996 Edition
ISO/IEC 9945-1: 1996
IEEE STANDARD OFFICE
ISBN 1-55937-573-6

Programming with POSIX threads
David R. Butenhof
Addison Wesley Longman, Inc., 1997
ISBN 0-201-63392-2

Programming

Advanced Programming in the UNIX Environment

Richard W. Stevens

Addison-Wesley, 1992

Debugging with GDB: The GNU Source-Level Debugger for GDB version 4.18

Richard Stallman, Cygnus Support

Out of Print

See <http://www.redhat.com/support/manuals/gnupro99r1/> for the online version

Open Source Development with CVS

Karl Franz Fogel

Coriolis Group, 1999, ISBN: 1576104907

Porting UNIX Software: From Download to Debug

Greg Lehey

O'Reilly, 1995, ISBN 1-56592-126-7

The Standard C Library

P.J. Plauger

Prentice Hall, 1992, ISBN 0-13-131509-9

Websites - Online Resources

Apple Computer's developer website (<http://www.apple.com/developer>) is a general repository for developer documentation. Additionally, the following sites provide more domain-specific information.

Apple's Public Source projects and Darwin OS

<http://www.publicsource.apple.com>

The Berkeley Software Distribution (BSD)

<http://www.FreeBSD.org>

<http://www.NetBSD.org>

<http://www.OpenBSD.org>

BSD Networking

<http://www.kohala.com/start/>

CVS (Concurrent Versions System)

<http://www.publicsource.apple.com/tools/cvs/cederquist>

Embedded C++

<http://www.caravan.net/ec2plus>

GDB, GNUPro Toolkit 99r1 Documentation

<http://www.redhat.com/support/manuals/gnupro99r1/>

The Internet Engineering Task Force (IETF)

<http://www.ietf.org>

jam

<http://www.perforce.com/jam/jam.html>

The PowerPC CPU

<http://www.motorola.com/SPS/PowerPC/>

The Single UNIX Specification Version 2

<http://www.opengroup.org/onlinepubs/007908799>

Stackable File Systems

http://www.isi.edu/~johnh/WORK/stacking_faq.htm

The USENIX Association, USENIX Proceedings

<http://www.usenix.org>

<http://www.usenix.org/publications/library/>
